

# AMS 250: An Introduction to High Performance Computing

## Parallel Programming Patterns Overview



**Shawfeng Dong**

[shaw@ucsc.edu](mailto:shaw@ucsc.edu)

(831) 502-7743

Applied Mathematics & Statistics  
University of California, Santa Cruz

# Outline

- Parallel programming models
- Dependencies
- Structured programming patterns overview

# Sequential Models

- von Neumann model

- Processing Unit, containing:


- Arithmetic Logic Unit, with processor registers
    - Control Unit, with instruction registers and program counter

- Memory, which stores both data and instructions

- External mass storage and I/O mechanism

- Stored-program computer

- CPU fetches instructions from memory, reads data from memory, decodes and executes instructions sequentially, then writes data back to memory



- Harvard model

- one dedicated set of address and data buses for reading data from and writing data to memory
  - another set of address and data buses for fetching instructions


# Parallel Models 101

- A parallel computer is simply a collection of *processors interconnected* in some manner to *coordinate* activities and *exchange data*
- Parallel models are those theoretical models that can be used as general frameworks for describing and analyzing parallel algorithms
  - ***Simplicity***: description, analysis, architecture independence
  - ***Implementability***: able to be realized, reflect performance
- Three common parallel models
  - **Directed acyclic graphs, shared-memory, network**

# Directed Acyclic Graphs (DAG)


- Captures data flow parallelism
- Nodes represent operations to be performed
  - Inputs are nodes with no incoming arcs
  - Output are nodes with no outgoing arcs
  - Think of nodes as tasks
- Arcs are paths for flow of data results
- DAG represents the operations of the algorithm and implies precedent constraints on their order

```
for (i=1; i<100; i++)  
    a[i] = a[i-1] + 100;
```




# Shared Memory Model

- Parallel extension of RAM model (PRAM)
  - Memory size is infinite
  - Number of processors is unbounded
  - Processors communicate via the memory
  - Every processor accesses any memory location in 1 cycle
  - Synchronous
    - All processors execute same algorithm synchronously
      - READ phase
      - COMPUTE phase
      - WRITE phase
    - Some subset of the processors can stay idle
  - Asynchronous



# Network Model

- $G = (N, E)$ 
  - $N$  are processing nodes
  - $E$  are bidirectional communication links
- Each processor node has its own memory
- No shared memory is available
- Network operation may be synchronous or asynchronous
- Requires communication primitives
  - Send ( $X, i$ )
  - Receive ( $Y, j$ )
- Captures **message passing** model for algorithm design



# Parallelism

- Ability to execute different parts of a computation concurrently on multiple processing elements
- Why do you want parallelism?
  - Shorter running time or handling more work
- What is being parallelized?
  - *Task*: instruction, statement, procedure, ...
  - *Data*: data flow, size, replication
  - Parallelism granularity
    - Coarse-grained versus fine-grained
- Thinking about parallelism
- Evaluation



# Parallel Algorithm

- Recipe to solve a problem “in parallel” on multiple processing elements
- Standard steps for constructing a parallel algorithm
  - Identify work that can be performed concurrently
  - Partition the concurrent work on separate processors
  - Properly manage input, output, and intermediate data
  - Coordinate data accesses and work to satisfy dependencies

# Parallelism Views

- Where can we find parallelism?
- Program (task) view
  - Statement level
    - Between program statements
    - Which statements can be executed at the same time?
  - Block level / Loop level / Routine level / Process level
    - Larger-grained program statements
- Data view
  - How is data operated on?
  - Where does data reside?
- Resource view

# Parallelism, Correctness, and Dependence

- Parallel execution, from any point of view, will be constrained by the sequence of operations needed to be performed for a correct result
- Parallel execution must address control, data, and system dependences
- A *dependency* arises when one operation depends on an earlier operation to complete and produce a result before this later operation can be performed
- We extend this notion of dependency to resources since some operations may depend on certain resources
  - For example, due to where data is located

# Executing Two Statements in Parallel

- Want to execute two statements in parallel
- On one processor:
  - statement 1;
  - statement 2;
- On two processors:

Processor 1:	Processor 2:
statement 1;	statement 2;
- Fundamental (*concurrent*) execution assumption
  - Processors execute independently of each other
  - No assumption made about speed of processor execution

# Sequential Consistency in Parallel Execution

- Case 1:

Processor 1:  
statement 1;

Processor 2:  
statement 2;

time  
↓

- Case 2:

Processor 1:  
statement 1;

Processor 2:  
statement 2;

time  
↓

- Sequential consistency

- Statements execution does not interfere with each other
- Computation results are the same (independent of order)

# Independent versus Dependent

- In other words the execution of  
    statement1;  
    statement2;  
must be equivalent to  
    statement2;  
    statement1;
- Their order of execution must not matter!
- If true, the statements are *independent* of each other
- Two statements are *dependent* when the order of their execution affects the computation outcome

# Examples

- Example 1

S1: a=1;

S2: b=1;

- Example 2

S1: a=1;

S2: b=a;

- Example 3

S1: a=f(x);

S2: a=b;

- Example 4

S1: a=b;

S2: b=1;

- Statements are independent

- Dependent (*true (flow) dependence*)

- Second is dependent on first

- Can you remove dependency?

- Dependent (*output dependence*)

- Second is dependent on first

- Can you remove dependency? How?

- Dependent (*anti-dependence*)

- First is dependent on second

- Can you remove dependency? How?

# True Dependence and Anti-Dependence

- Given statements S1 and S2,  
S1;  
S2;
- S2 has a *true (flow) dependence* on S1  
if and only if  
S2 reads a value written by S1
- S2 has a *anti-dependence* on S1  
if and only if  
S2 writes a value read by S1


$$\begin{array}{c} X \\ \vdots \\ = X \end{array} \left. \vphantom{\begin{array}{c} X \\ \vdots \\ = X \end{array}} \right] \delta$$

$$\begin{array}{c} = X \\ \vdots \\ X = \end{array} \left. \vphantom{\begin{array}{c} = X \\ \vdots \\ X = \end{array}} \right] \delta^{-1}$$



# Output Dependence

- Given statements S1 and S2,  
    S1;  
    S2;
- S2 has an *output dependence* on S1  
    if and only if  
    S2 writes a variable written by S1



- Anti- and output dependences are “name” dependencies
  - Are they “true” dependences?
- How can you get rid of output dependences?
  - Are there cases where you can not?

# Statement Dependency Graphs

- We can use graphs to show dependence relationships


- Example

S1: a=1;

S2: b=a;

S3: a=b+1;

S4: c=a;



- $S_2 \delta S_3$  :  $S_3$  is flow-dependent on  $S_2$
- $S_1 \delta^0 S_3$  :  $S_3$  is output-dependent on  $S_1$
- $S_2 \delta^{-1} S_3$  :  $S_3$  is anti-dependent on  $S_2$

# When can two statements execute in parallel?

- Statements S1 and S2 can execute in parallel if and only if there are *no dependences* between S1 and S2
  - True dependences
  - Anti-dependences
  - Output dependences
- Some dependences can be removed by modifying the program
  - Rearranging statements
  - Eliminating statements

# How do you compute dependence?

- Data dependence relations can be found by comparing the IN and OUT sets of each node
- The IN and OUT sets of a statement  $S$  are defined as:
  - $IN(S)$  : set of memory locations (variables) that may be used in  $S$
  - $OUT(S)$  : set of memory locations (variables) that may be modified by  $S$
- Note that these sets include all memory locations that may be fetched or modified
- As such, the sets can be conservatively large

# IN / OUT Sets and Computing Dependence

- Assuming that there is a path from **S1** to **S2** , the following shows how to intersect the IN and OUT sets to test for data dependence

$out(S_1) \cap in(S_2) \neq \emptyset$        $S_1 \delta S_2$       flow dependence

$in(S_1) \cap out(S_2) \neq \emptyset$        $S_1 \delta^{-1} S_2$       anti - dependence

$out(S_1) \cap out(S_2) \neq \emptyset$        $S_1 \delta^0 S_2$       output dependence

# Loop-Level Parallelism

- Significant parallelism can be identified within loops

```
for (i=0; i<100; i++)  
    S1: a[i] = i;  
  
for (i=0; i<100; i++) {  
    S1: a[i] = i;  
    S2: b[i] = 2*i;  
}
```

- Dependencies? What about  $i$ , the loop index?
- *DOALL* loop (a.k.a. *foreach* loop)
  - All iterations are independent of each other
  - All statements can be executed in parallel at the same time


# Iteration Space

- Unroll loop into separate statements / iterations
- Show dependences between iterations

```
for (i=0; i<100; i++)  
  s1: a[i] = i;
```




```
for (i=0; i<100; i++) {  
  s1: a[i] = i;  
  s2: b[i] = 2*i;  
}
```



# Multi-Loop Parallelism

- Significant parallelism can be identified between loops

```
for (i=0; i<100; i++) a[i] = i;  
for (i=0; i<100; i++) b[i] = i;
```




- Dependencies?
- How much parallelism is available?
- Given 4 processors, how much parallelism is possible?
- What parallelism is achievable with 50 processors?



# Loops with Dependencies

Case 1:


```
for (i=1; i<100; i++)  
    a[i] = a[i-1] + 100;
```



- Dependencies?
  - What type?
- Is the Case 1 loop parallelizable?
- Is the Case 2 loop parallelizable?

Case 2:

```
for (i=5; i<100; i++)  
    a[i-5] = a[i] + 100;
```



# Another Loop Example

```
for (i=1; i<100; i++)  
    a[i] = f(a[i-1]);
```

- Dependencies?
  - What type?
- Loop iterations are not parallelizable
  - Why not?

# Loop Dependencies

- A *loop-carried* dependence is a dependence that is present only if the statements are part of the execution of a loop (i.e., between two statements instances in two different iterations of a loop)
- Otherwise, it is *loop-independent*, including between two statements instances in the same loop iteration
- Loop-carried dependences can prevent loop iteration parallelization
- The dependence is *lexically forward* if the source comes before the target or *lexically backward* otherwise
  - Unroll the loop to see


# Loop Dependence Example

```
for (i=0; i<100; i++)  
    a[i+10] = f(a[i]);
```

- Dependencies?
  - Between a[10], a[20], ...
  - Between a[11], a[21], ...
- Some parallel execution is possible
  - How much?

# Dependencies Between Iterations

```
for (i=1; i<100; i++) {  
    S1: a[i] = ...;  
    S2: ... = a[i-1];  
}
```



- Dependencies?
  - Between  $a[i]$  and  $a[i-1]$
- Is parallelism possible?
  - Statements can be executed in “pipeline” manner

# Another Loop Dependence Example

```
for (i=0; i<100; i++)  
    for (j=1; j<100; j++)  
        a[i][j] = f(a[i][j-1]);
```

- Dependencies?
  - Loop-independent dependence on i
  - Loop-carried dependence on j
- Which loop can be parallelized?
  - Outer loop parallelizable
  - Inner loop cannot be parallelized

# Still Another Loop Dependence Example

```
for (j=1; j<100; j++)  
    for (i=0; i<100; i++)  
        a[i][j] = f(a[i][j-1]);
```

- Dependencies?
  - Loop-independent dependence on i
  - Loop-carried dependence on j
- Which loop can be parallelized?
  - Inner loop parallelizable
  - Outer loop cannot be parallelized
  - Less desirable (why?)

# Key Ideas for Dependency Analysis

- To execute in parallel:
  - Statement order must not matter
  - Statements must not have dependences
- Some dependences can be removed
- Some dependences may not be obvious




# Dependencies and Synchronization

- How is parallelism achieved when having dependencies?
  - Think about concurrency
  - Some parts of the execution are independent
  - Some parts of the execution are dependent
- Must control ordering of events on different processors (cores)
  - Dependencies pose constraints on parallel event ordering
  - Partial ordering of execution action
- Use synchronization mechanisms
  - Need for concurrent execution too
  - Maintains partial order

# Structured Programming with Patterns

- Patterns are “best practices” for solving specific problems.
- Patterns can be used to organize your code, leading to algorithms that are more scalable and maintainable.
- A pattern supports a particular “algorithmic structure” with an efficient implementation.
- Good parallel programming models support a set of useful parallel patterns with low-overhead implementations.



Graphical notation for the fundamental components of algorithms

# Structured Serial Patterns

The following patterns are the basis of “**structured programming**” for serial computation:

- Sequence
- Selection
- Iteration
- Nesting
- Functions
- Recursion
- Random read
- Random write
- Stack allocation
- Heap allocation
- Objects
- Closures

*Using these patterns, “goto” can (mostly) be eliminated and the maintainability of software improved.*

# Structured Parallel Patterns

The following additional parallel patterns can be used for “**structured parallel programming**”:

- Superscalar sequence
- Speculative selection
- Map
- Recurrence
- Scan
- Reduce
- Pack/expand
- Fork/join
- Pipeline
- Partition
- Segmentation
- Stencil
- Search/match
- Gather
- Merge scatter
- Priority scatter
- Permutation scatter
- Atomic scatter

*Using these patterns, threads and vector intrinsics can (mostly) be eliminated and the maintainability of software improved.*

# Some Basic Patterns

- **Serial:** Sequence
  - **Parallel:** Superscalar Sequence
- **Serial:** Iteration
  - **Parallel:** Map, Reduction, Scan, Recurrence...

# (Serial) Sequence



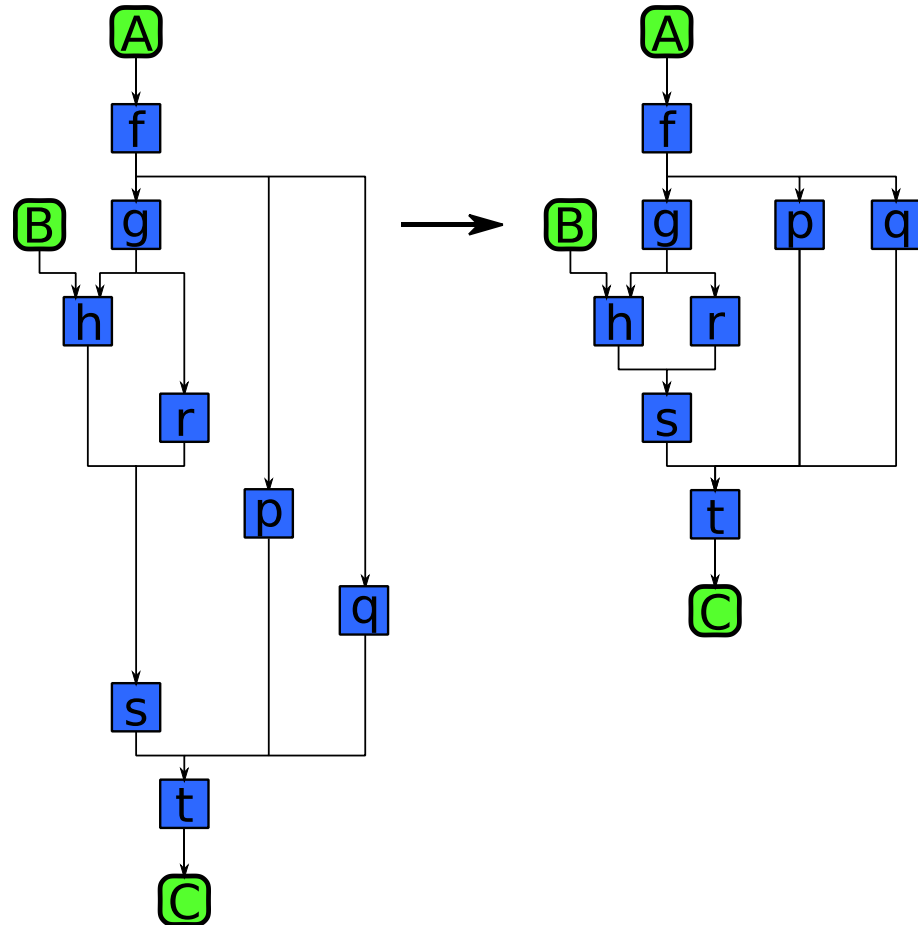
A serial sequence is executed in the exact order given:

$F = f(A);$

$G = g(F);$

$B = h(G);$

# Superscalar Sequence




Developer writes “serial” code:

```
F = f(A);  
G = g(F);  
H = h(B,G);  
R = r(G);  
P = p(F);  
Q = q(F);  
S = s(H,R);  
C = t(S,P,Q);
```

- Tasks ordered only by data dependencies
- Tasks can run whenever input data is ready

# (Serial) Iteration



The iteration pattern repeats some section of code as long as a condition holds


```
while (c) {  
    f();  
}
```

Each iteration can depend on values computed in any earlier iteration.

The loop can be terminated at any point based on computations in any iteration



# (Serial) Countable Iteration



The iteration pattern repeats some section of code a specific number of times

```
for (i = 0; i < n; ++i) {  
    f();  
}
```


This is the same as

```
i = 0;  
while (i < n) {  
    f();  
    ++i;  
}
```

# Parallel “Iteration”

- The serial iteration pattern actually maps to several *different* parallel patterns
- It depends on whether and how iterations depend on each other...
- Most parallel patterns arising from iteration require a fixed number of invocations of the body, known in advance

# Map




**Examples:** gamma correction and thresholding in images; color space conversions; Monte Carlo sampling; ray tracing.

- *Map* replicates a function over every element of an index set
- The index set may be abstract or associated with the elements of an array.

```
for (i=0; i<n; ++i) {  
    f(A[i]);  
}
```

- Map replaces *one specific* usage of iteration in serial programs: *independent operations*.

# Reduction




**Examples:** averaging of Monte Carlo samples; convergence testing; image comparison metrics; matrix operations.

- *Reduction* combines every element in a collection into one element using an associative operator.

```
b = 0;  
for (i=0; i<n; ++i) {  
    b += f(B[i]);  
}
```

- Reordering of the operations is often needed to allow for parallelism.
- A tree reordering requires associativity.

# Scan




**Examples:** random number generation, pack, tabulated integration, time series analysis

- *Scan* computes all partial reductions of a collection

```
A[0] = B[0] + init;  
for (i=1; i<n; ++i) {  
    A[i] = B[i] + A[i-1];  
}
```

- Operator must be (at least) associative.
- Diagram shows one possible parallel implementation using three-phase strategy


# Geometric Decomposition/Partition



**Examples:** JPG and other macroblock compression; divide-and-conquer matrix multiplication; coherency optimization for cone-beam recon.

- *Geometric decomposition* breaks an input collection into sub-collections
- *Partition* is a special case where sub-collections do not overlap
- Does not move data, it just provides an alternative “view” of its organization


# Stencil



**Examples:** signal filtering including convolution, median, anisotropic diffusion

- *Stencil* applies a function to neighbourhoods of a collection.
- Neighbourhoods are given by set of relative offsets.
- Boundary conditions need to be considered, but majority of computation is in interior.

# nD Stencil





**Examples:** image filtering including convolution, median, anisotropic diffusion; simulation including fluid flow, electromagnetic, and financial PDE solvers, lattice QCD

- *nD Stencil* applies a function to neighbourhoods of an nD array
- Neighbourhoods are given by set of relative offsets
- Boundary conditions need to be considered

```
for (int i = 1, i < N; i++)  
  for (int j = 1, j < M; j++)  
    a_new[i][j] = 0.25 *  
      (a[i-1][j] +  
       a[i+1][j] +  
       a[i][j-1] +  
       a[i][j+1]);
```




# Implementing Stencil



**Vectorization** can include converting regular reads into a set of shifts.

**Strip-mining** reuses previously read inputs within serialized chunks.

# Recurrence




- *Recurrence* results from loop nests with both input and output dependencies between iterations
- Can also result from iterated stencils


**Examples:** Simulation including fluid flow, electromagnetic, and financial PDE solvers, lattice QCD, sequence alignment and pattern matching

# Recurrence

```
for (int i = 1; i < N; i++) {  
    for (int j = 1; j < M; j++) {  
        A[i][j] = f(  
            A[i-1][j],  
            A[i][j-1],  
            A[i-1][j-1],  
            A[i][j]);  
    }  
}
```




# Recurrence Hyperplane Sweep




- Multidimensional recurrences can *always* be parallelized
- Leslie Lamport's hyperplane separation theorem:
  - Choose hyperplane with inputs and outputs on opposite sides
  - Sweep through data perpendicular to hyperplane

# Rotated Recurrence




- Rotate recurrence to see sweep more clearly

# Tiled Recurrence




- Can partition recurrence to get a better compute vs. bandwidth ratio

# Tiled Recurrence



- Remove all non-redundant data dependencies


# Recursively Tiled Recurrences



- Rotate back: same recurrence at a different scale!
- Leads to recursive cache-oblivious divide-and-conquer algorithm
- Implement with fork-join.




# Pipeline



- *Pipeline* uses a sequence of stages that transform a flow of data
- Some stages may retain state
- Pipeline connects tasks in a producer-consumer manner

**Examples:** image filtering, data compression and decompression, signal processing

# Pipeline




- Parallelize pipeline by
  - Running different stages in parallel
  - Running *multiple copies* of stateless stages in parallel
- Running multiple copies of stateless stages in parallel requires reordering of outputs
- Need to manage buffering between stages

# Recursive Patterns


- Recursion is an important “universal” serial pattern
  - Recursion leads to functional programming
  - Iteration leads to procedural programming
- Structural recursion: nesting of components
- Dynamic recursion: nesting of behaviors

# Nesting: Recursive Composition



**Nesting Pattern:** A compositional pattern. Nesting allows other patterns to be composed in a hierarchy so that any task block in the above diagram can be replaced with a pattern with the same input/output and dependencies.

# Fork-Join: Efficient Nesting



- Fork-join can be nested
- Spreads cost of work distribution and synchronization.


Recursive fork-join enables high parallelism.

# Other Parallel Patterns


- **Futures:** similar to fork-join, but tasks do not need to be nested hierarchically
- **Speculative Selection:** general version of serial selection where the condition and both outcomes can all run in parallel
- **Workpile:** general map pattern where each instance of elemental function can generate more instances, adding to the “pile” of work
- **Search:** finds some data in a collection that meets some criteria
- **Segmentation:** operations on subdivided, non-overlapping, non-uniformly sized partitions of 1D collections
- **Expand:** a combination of pack and map
- **Category Reduction:** Given a collection of elements each with a label, find all elements with same label and reduce them

# Parallel Patterns: Overview


Superscalar sequence




Speculative selection




Map




Stencil




Fork-Join




Pipeline




Geometric decomposition




Partition




Gather




Scatter




Category Reduction




Recurrence




Reduction




Scan




Expand



Pack



Split



# Semantics and Implementation

## Semantics: *What*

- The intended meaning as seen from the “outside”
- For example, for scan: compute all partial reductions given an associative operator

## Implementation: *How*

- How it executes in practice, as seen from the “inside”
- For example, for scan: partition, serial reduction in each partition, scan of reductions, serial scan in each partition.
- *Many implementations may be possible*
- Parallelization may require reordering of operations
- Patterns should not over-constrain the ordering; only the important ordering constraints are specified in the semantics
- Patterns may also specify additional constraints, i.e. associativity of operators



# POSIX Threads

- POSIX standard multi-threading interface
  - For general multi-threaded concurrent programming
  - Defined as a set of **C** programming language types, functions and constants
  - Largely independent across implementations, and broadly supported
  - Common target for library and language implementation
- Provides primitives for
  - Thread management - creating, joining threads etc.
  - Synchronization
- POSIX Threads (pthreads) specification:  
<http://pubs.opengroup.org/onlinepubs/9699919799/basedefs/pthread.h.html>
- POSIX Threads Programming tutorial:  
<https://computing.llnwd.net/tutorials/pthreads/>

# C++11 Multithreading

- C++11 standardizes support for multithreaded programming:
  - a memory model which allows multiple threads to co-exist in a program
  - library support for interaction between threads
- C++11 standard library includes:
  - Atomics
  - Threads
  - Mutexes
  - Conditional Variables
  - Futures and Promises
- C++11 multithreading reference:  
<http://www.cplusplus.com/reference/multithreading/>
- A good book on C++11 multithreading: **C++ Concurrency in Action: Practical Multithreading**, by Anthony Williams, Manning Publications, 2012

# Grand Central Dispatch

- Developed by Apple
- Available on OS X 10.6 and later, iOS 4 and later
  - Open-sourced under the Apache license: <https://libdispatch.macosforge.org>
- An implementation of **task parallelism** based on the **thread pool** pattern
  - Still uses threads at the low level but abstracts them away from the programmer
  - Allows tasks to be queued, then schedules them to execute on any of the available processor cores
  - A task can be expressed either as a function or as a “block”
  - Grand Central Dispatch (GCD) reference:  
[https://developer.apple.com/library/ios/documentation/Performance/Reference/GCD\\_libdispatch\\_Ref/](https://developer.apple.com/library/ios/documentation/Performance/Reference/GCD_libdispatch_Ref/)

# Thread Building Blocks

- Threading Building Blocks (TBB) is a **C++** template library developed by Intel for parallel programming on multi-core processors
  - A TBB program specifies graphs of dependent tasks according to *algorithms/patterns*, instead of manipulating threads
  - TBB implements *work stealing* to balance a parallel workload across available processing cores in order to increase core utilization and therefore scaling
  - TBB includes efficient low-level primitives (atomics, memory allocation, etc.)
- TBB is available:
  - both commercially as a binary distribution with support
  - and as open-source software: <https://www.threadingbuildingblocks.org/>
- TBB tutorial: <https://www.threadingbuildingblocks.org/intel-tbb-tutorial>

# TBB 4.0 Components

## Parallel Algorithms

`parallel_for`  
`parallel_for_each`  
`parallel_invoke`  
`parallel_do`  
`parallel_scan`  
`parallel_sort`  
`parallel_[deterministic]_reduce`

## Macro Dataflow

`parallel_pipeline`  
`tbb::flow::...`

## Concurrent Containers

`concurrent_hash_map`  
`concurrent_unordered_{map,set}`  
`concurrent_[bounded_]queue`  
`concurrent_priority_queue`  
`concurrent_vector`

## Task scheduler

`task_group, structured_task_group`  
`task`  
`task_scheduler_init`  
`task_scheduler_observer`

## Synchronization Primitives

`atomic, condition_variable`  
`[recursive_]mutex`  
`{spin,queuing,null}[_rw]_mutex`  
`critical_section, reader_writer_lock`

## Threads

`std::thread`

## Thread Local Storage

`combinable`  
`enumerable_thread_specific`

## Memory Allocation

`tbb_allocator`  
`zero_allocator`  
`cache_aligned_allocator`  
`scalable_allocator`

# Intel Cilk Plus

- Intel Cilk Plus is an *extension* to **C** and **C++** that simplifies the expression of task and data parallelism:
  - Fork-join parallel programming model
  - Serial semantics if keywords are ignored (serial elision)
  - Efficient work-stealing load balancing
  - Supports vector parallelism via array slices and elemental functions
- Cilk Plus is available:
  - both commercially as a binary distribution with support
  - and as open-source software: <https://www.cilkplus.org/>
- Cilk Plus tutorial: <https://www.cilkplus.org/cilk-plus-tutorial>

# Summary of Intel Cilk Plus

## Thread Parallelism

cilk\_spawn  
cilk\_sync  
cilk\_for

## Vector Parallelism

array notation  
#pragma simd  
elemental functions

## Reducers

reducer  
reducer\_op{add,and,or,xor}  
reducer\_{min,max}{\_index}  
reducer\_list\_{append,prepend}  
reducer\_ostream  
reducer\_string  
holder