

AMS 250: An Introduction to High Performance Computing

Advanced MPI



Shawfeng Dong

shaw@ucsc.edu

(831) 502-7743

Astronomy & Astrophysics

University of California, Santa Cruz

Outline


- MPI + OpenMP
- One-Sided Communications with MPI

MPI + X

- For the past 20+ years, HPC programming is dominated by
 - either a flat model, with **MPI** across nodes as well as cores within a node
 - or a hybrid model, with **MPI** across the nodes and **OpenMP** shared memory parallelism across the cores in a node
- The promise of **PGAS** (Partitioned Global Address Space) has failed to materialize – more on PGAS later
- MPI continues to be enhanced and evolved – MPI 3.1 released in June 2015
- In the exascale era, the programming model will likely be MPI + X, where X is one of:
 - OpenMP 3 (if accelerators don't persist)
 - OpenMP 4 with *target* directives
 - OpenACC
 - CUDA or OpenCL
 - PGAS

MPI + OpenMP

- Modern HPC systems are predominantly clusters of SMPs, and increasing heterogeneous:



- Flat MPI model:
 - ☺ simpler programming model: only one level of parallelism and only one API
 - ☹ doesn't take advantage of the shared data across the ranks on the same node, requiring message and buffer management across all ranks
- Hybrid MPI + OpenMP model:
 - ☹ more complex programming model
 - ☺ maps nicely to today's prevalent architecture

MPI and Threads

- MPI libraries can vary in their level of thread support:
 - **MPI_THREAD_SINGLE** – Level 0: Only one thread will execute.
 - **MPI_THREAD_FUNNELED** – Level 1: The process may be multi-threaded, but only the main thread will make MPI calls (all MPI calls are funneled to the main thread).
 - **MPI_THREAD_SERIALIZED** – Level 2: The process may be multi-threaded, and multiple threads may make MPI calls, but only one at a time. That is, calls are not made concurrently from two distinct threads as all MPI calls are serialized.
 - **MPI_THREAD_MULTIPLE** – Level 3: Multiple threads may make MPI calls with no restrictions.
- OpenMP and Pthreads are common models for thread parallelism

MPI and Threads (cont'd)

- An implementation is not required to support levels higher than **MPI_THREAD_SINGLE**; that is, an implementation is not required to be thread safe
- A fully thread-compliant implementation will support **MPI_THREAD_MULTIPLE**
- Call **MPI_Init_thread** (instead of **MPI_Init**) to initialize the multi-threaded MPI execution environment

C	<code>int MPI_Init_thread(int *argc, char ***argv, int required, int *provided)</code>
Fortran	<code>MPI_INIT_THREAD(REQUIRED, PROVIDED, IERROR) INTEGER REQUIRED, PROVIDED, IERROR</code>

- A portable program that does not call **MPI_Init_thread** should assume that only **MPI_THREAD_SINGLE** is supported

An example to query the level of thread support

```
#include "mpi.h"
#include <stdio.h>

int main(int argc, char* argv[]) {
    int rank;
    int thread_level, thread_is_main;
    /* MPI_Init(&argc, &argv) */
    MPI_Init_thread(&argc, &argv, MPI_THREAD_MULTIPLE, &thread_level);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank == 0) {
        printf("MPI_THREAD_SINGLE = %d\n", MPI_THREAD_SINGLE);
        printf("MPI_THREAD_FUNNELED = %d\n", MPI_THREAD_FUNNELED);
        printf("MPI_THREAD_SERIALIZED = %d\n", MPI_THREAD_SERIALIZED);
        printf("MPI_THREAD_MULTIPLE = %d\n", MPI_THREAD_MULTIPLE);
        printf("I asked for level 3, and I got level %d\n", thread_level);
        MPI_Query_thread( &thread_level );
        MPI_Is_thread_main( &thread_is_main );
        printf("Thread level is %d\n", thread_level);
        printf("thread_is_main = %d\n", thread_is_main);
    }
    MPI_Finalize();
    return 0;
}
```

Try it out on Cori

```
shawdong@cori06:> cc mpi_thread.c -o mpi_thread.x
shawdong@cori06:> salloc -N 1 -p debug -L SCRATCH -C haswell

shawdong@nid12954:> srun -n 2 ./mpi_thread.x
MPI_THREAD_SINGLE = 0
MPI_THREAD_FUNNELED = 1
MPI_THREAD_SERIALIZED = 2
MPI_THREAD_MULTIPLE = 3
I asked for level 3, and I got level 2
Thread level is 2
thread_is_main = 1
```

Cray MPI only supports up to MPI_THREAD_SERIALIZED (level 2)!

Try it out on Hyades (Intel MPI)

```
$ mpiicc -mt_mpi mpi_thread.c -o mpi_thread.x
$ mpirun -n 2 ./mpi_thread.x
MPI_THREAD_SINGLE = 0
MPI_THREAD_FUNNELED = 1
MPI_THREAD_SERIALIZED = 2
MPI_THREAD_MULTIPLE = 3
I asked for level 3, and I got level 3
Thread level is 3
thread_is_main = 1
```

Note:

- You must link with the thread safe version of the Intel MPI library (-mt_mpi)
- Intel MPI offers all 3 levels of thread support

MPI_THREAD_MULTIPLE

- Multiple threads may make MPI calls with no restrictions.
- When multiple threads make MPI calls concurrently, the outcome will be as if the calls executed sequentially in some (any) order
- Blocking MPI calls will block only the calling thread and will not prevent other threads from running or executing MPI functions
- It is the user's responsibility to prevent races when threads in the same application post conflicting MPI calls
- User must ensure that collective operations on the same communicator, window, or file handle are correctly ordered among threads
- “Thread-safe” usually means `MPI_THREAD_MULTIPLE`

MPI_THREAD_SERIALIZED

Multiple threads may make MPI calls, but only one at a time.

```
#pragma omp parallel
...
#pragma omp atomic
{
... MPI calls allowed here ...
}
```

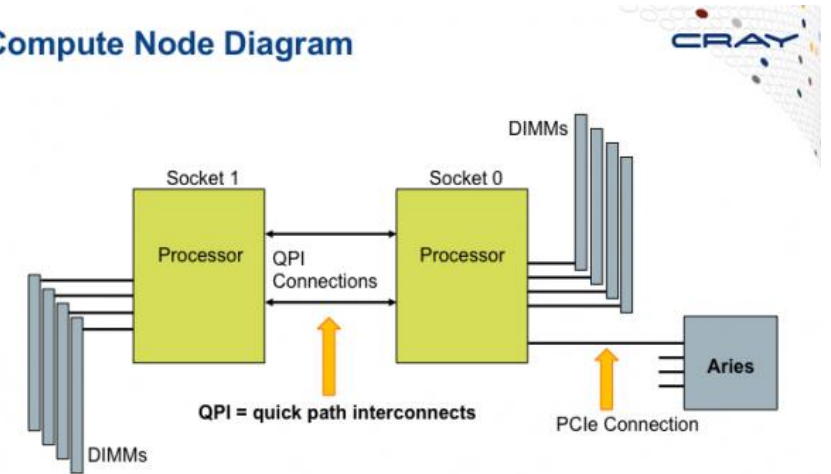
MPI_THREAD_FUNNELED

- All of the MPI calls are made by the master thread. That is, all MPI calls are:
 - either outside OpenMP parallel regions
 - or inside OpenMP master regions
 - or guarded by an **MPI_Is_thread_main** call (same thread that's called **MPI_Init_thread**)
- Probably sufficient for typical MPI + OpenMP hybrid programs

Edison Compute Nodes

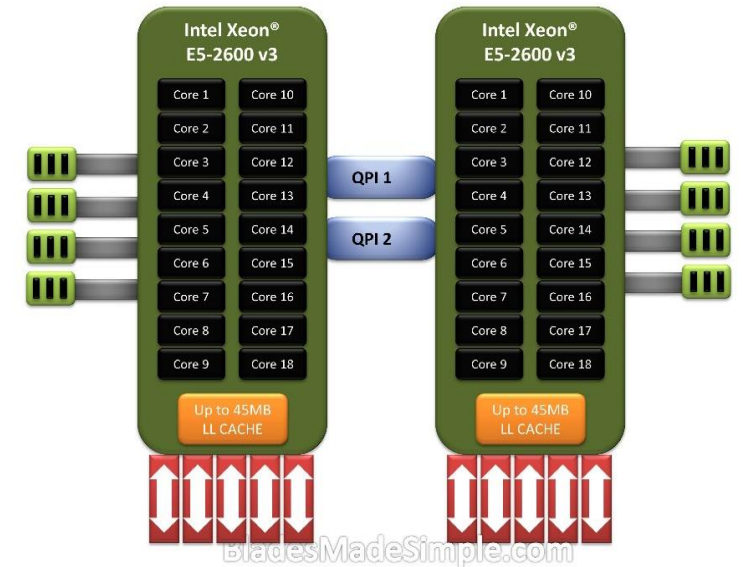
- Two Intel "Ivy Bridge" Xeon E5-2695 v2 processors per node
 - 12 cores per processor
 - 256-bit/cycle ring bus interconnect between on-die cores
 - integrated memory controller and PCIe controller
 - 2 QPI links between the 2 processors
- Two NUMA nodes per node
 - one NUMA node per processor/socket
 - local memory bandwidth: 59.7 GB/s
 - remote memory bandwidth: 32.0 GB/s
- An Aries Network Interface Controller (NIC)
 - 0.25 μ s to 3.7 μ s MPI latency
 - ~8GB/sec MPI bandwidth

Compute Node Diagram



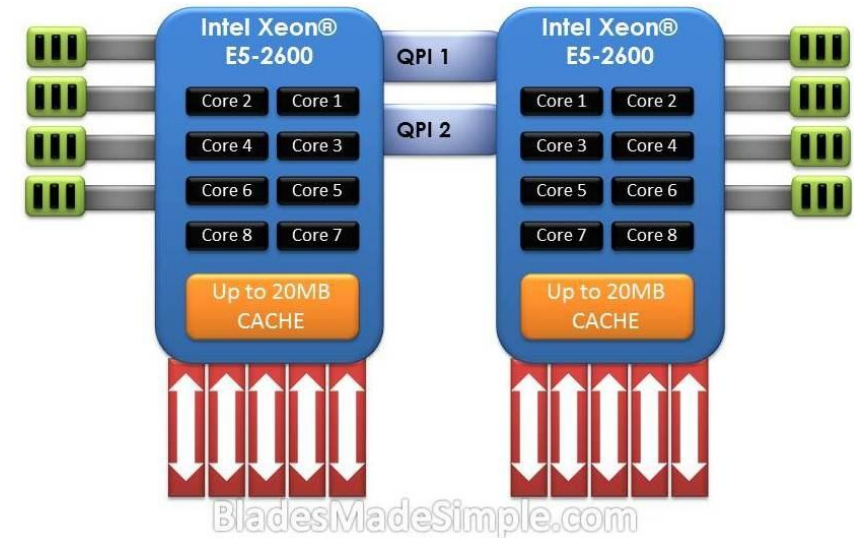
Cori Haswell Compute Nodes

- Two Intel "Haswell" Xeon E5-2698 v3 processors per node
 - 16 cores per processor
 - 256-bit/cycle ring bus interconnect between on-die cores
 - integrated memory controller and PCIe controller
 - 2 QPI links between the 2 processors
- Two NUMA nodes per node
 - one NUMA node per processor/socket
 - local memory bandwidth: 68.3 GB/s
 - remote memory bandwidth: 38.4 GB/s
- An Aries Network Interface Controller (NIC)
 - 0.25 μ s to 3.7 μ s MPI latency
 - ~8GB/sec MPI bandwidth



Hyades Compute Nodes

- Two Intel "Sandy Bridge" Xeon E5-2650 processors per node
 - 8 cores per processor
 - 256-bit/cycle ring bus interconnect between on-die cores
 - integrated memory controller and PCIe controller
 - 2 QPI links between the 2 processors
- Two NUMA nodes per node
 - one NUMA node per processor/socket
 - local memory bandwidth: 51.2 GB/s
 - remote memory bandwidth: 32.0 GB/s
- Mellanox ConnectX-2 VPI QDR InfiniBand HCA
 - 5 GB/s signaling rate
 - 4 GB/s data rate
 - for MPI communications across nodes




Running Modes


Pure MPI

1 MPI Task
Thread on each Core


16 MPI Tasks



2 MPI Tasks
8 threads/task



1 MPI Tasks
16 threads/task



Master Thread of MPI Task

■ MPI Task on Core

■ Master Thread of MPI Task

■ Slave Thread of MPI Task

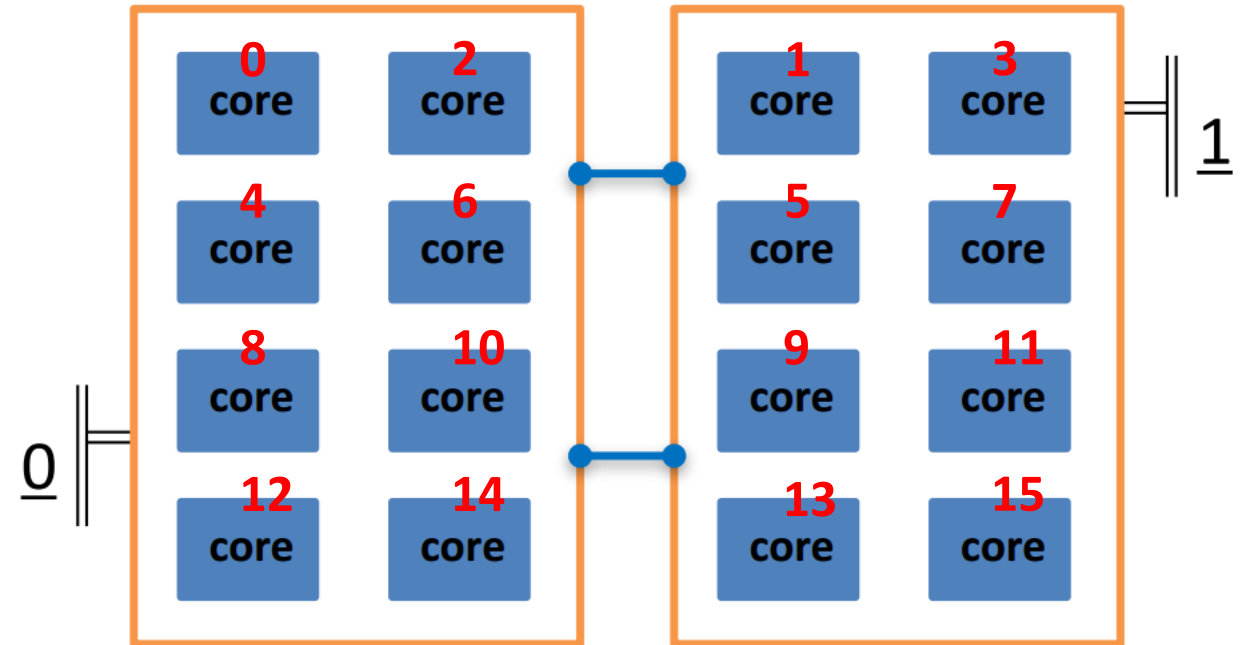
NUMA Operations

- Each process/thread is executed by a core and has access to a certain memory space
 - Core assigned by **process affinity**
 - Memory allocation assigned by **memory policy**
- The control of **process affinity** and **memory policy** using NUMA operations
 - NUMA Control is managed by the kernel (default).
 - Users can alter kernel policies by manually setting **process affinity** and **memory policy**, within a program through C/Fortran API, or using *numactl*.

Processor IDs

```
$ egrep 'processor|physical id|core id' /proc/cpuinfo
```

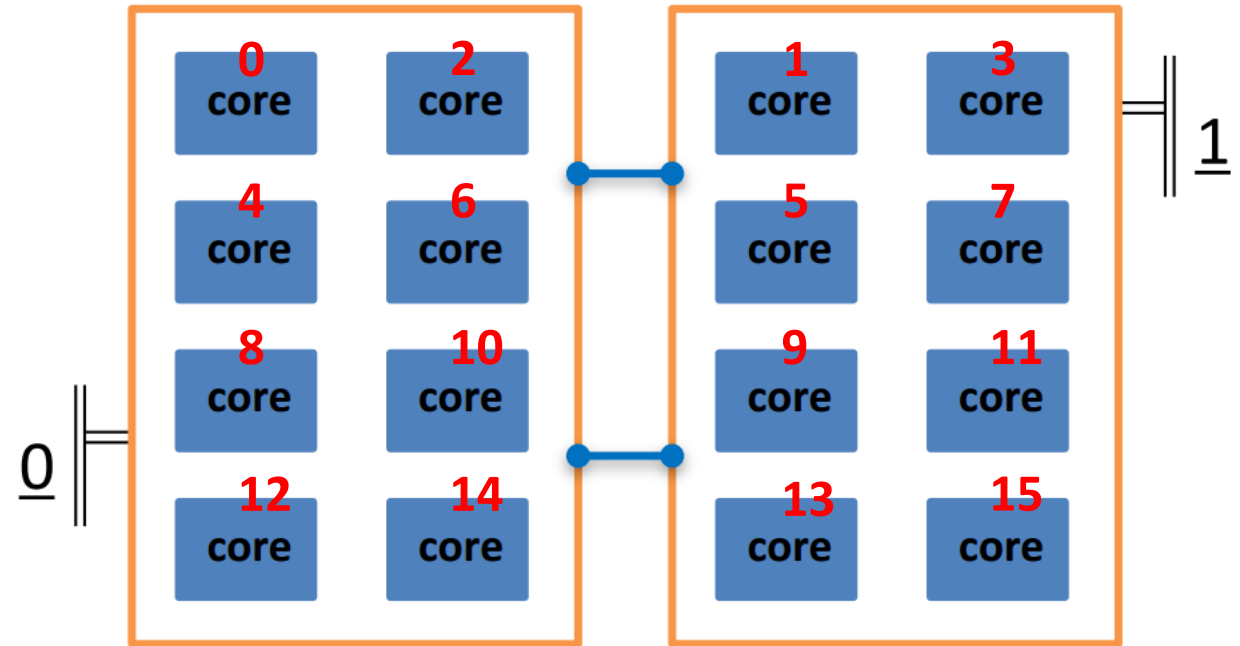
```
processor      : 0  
physical id   : 0  
core id       : 0  
processor      : 1  
physical id   : 1  
core id       : 0  
processor      : 2  
physical id   : 0  
core id       : 1  
processor      : 3  
physical id   : 1  
core id       : 1  
...
```



Hyades Compute Node

numactl

```
$ numactl --hardware
available: 2 nodes (0-1)
node 0 cpus: 0 2 4 6 8 10 12 14
node 0 size: 32722 MB
node 0 free: 24463 MB
node 1 cpus: 1 3 5 7 9 11 13 15
node 1 size: 32768 MB
node 1 free: 26229 MB
node distances:
node    0    1
  0:   10   20
  1:   20   10
```



Hyades Compute Node

Cori Haswell Compute Node

```
cori08:> salloc -N 1 -p debug -t 00:10:00 -C haswell
```

```
nid12954:> numactl -H
```

```
available: 2 nodes (0-1)
```

```
node 0 cpus: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15  
32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47
```

```
node 0 size: 64430 MB
```

```
node 0 free: 62803 MB
```

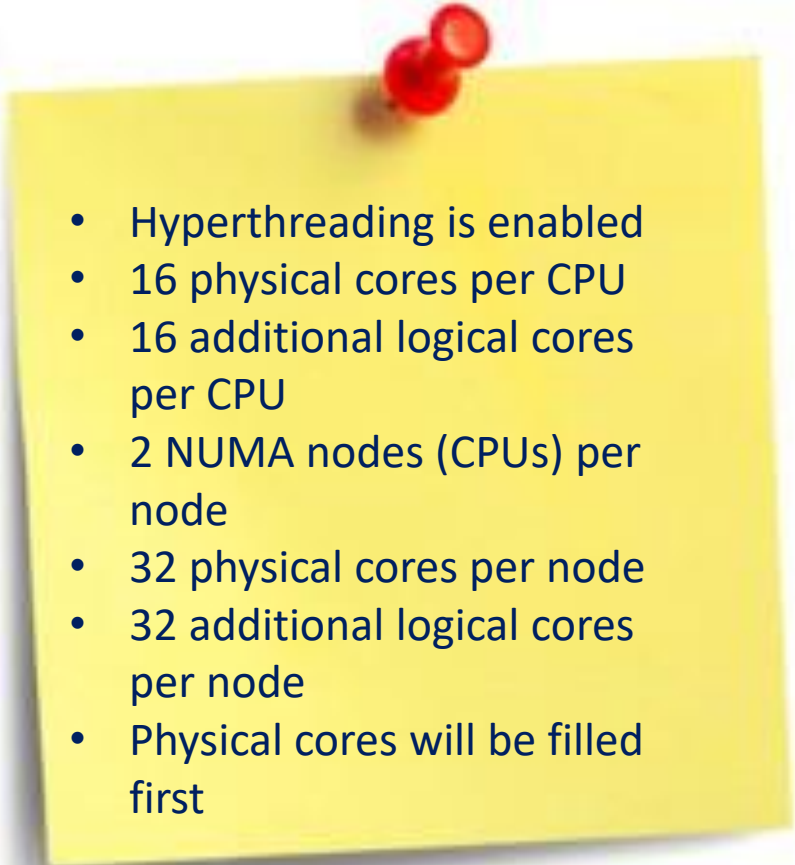
```
node 1 cpus: 16 17 18 19 20 21 22 23 24 25 26 27 28  
29 30 31 48 49 50 51 52 53 54 55 56 57 58 59 60 61  
62 63
```

```
node 1 size: 64635 MB
```

```
node 1 free: 63552 MB
```

```
node distances:
```

```
node    0    1  
  0:   10   21  
  1:   21   10
```

- 
- Hyperthreading is enabled
 - 16 physical cores per CPU
 - 16 additional logical cores per CPU
 - 2 NUMA nodes (CPUs) per node
 - 32 physical cores per node
 - 32 additional logical cores per node
 - Physical cores will be filled first

Cori KNL Compute Node: SNC4 + Cache

```
cori08:> salloc -N 1 -p debug -t 00:10:00 -L SCRATCH -C knl,snc4,cache
```

```
nid12253:~> numactl --hardware
```

```
available: 4 nodes (0-3)
```

```
node 0 cpus: 0-17,68-85,136-153,204-221
```

```
node 0 size: 24065 MB
```

```
node 1 cpus: 18-35,86-103,154-171,222-239
```

```
node 1 size: 24232 MB
```

```
node 2 cpus: 36-51,104-119,172-187,240-255
```

```
node 2 size: 24233 MB
```

```
node 3 cpus: 52-67,120-135,188-203,256-271
```

```
node 3 size: 24230 MB
```

```
node distances:
```

node	0	1	2	3
0:	10	21	21	21
1:	21	10	21	21
2:	21	21	10	21
3:	21	21	21	10

Cori KNL Compute Node: SNC4 + Flat

```
cori08:> salloc -N 1 -p debug -t 00:10:00 -L SCRATCH -C knl,snc4,flat
```

```
nid04893:~> numactl -H
```

```
available: 8 nodes (0-7)
```

```
node 0 cpus: 0-17,68-85,136-153,204-221
```

```
node 0 size: 24065 MB
```

```
node 1 cpus: 18-35,86-103,154-171,222-239
```

```
node 1 size: 24232 MB
```

```
node 2 cpus: 36-51,104-119,172-187,240-255
```

```
node 2 size: 24233 MB
```

```
node 3 cpus: 52-67,120-135,188-203,256-271
```

```
node 3 size: 24233 MB
```

```
node 4 cpus:
```

```
node 4 size: 4039 MB
```

```
node 5 cpus:
```

```
node 5 size: 4039 MB
```

```
node 6 cpus:
```

```
node 6 size: 4039 MB
```

```
node 7 cpus:
```

```
node 7 size: 4037 MB
```

node distances:

node	0	1	2	3	4	5	6	7
0:	10	21	21	21	31	41	41	41
1:	21	10	21	21	41	31	41	41
2:	21	21	10	21	41	41	31	41
3:	21	21	21	10	41	41	41	31
4:	31	41	41	41	10	41	41	41
5:	41	31	41	41	41	10	41	41
6:	41	41	31	41	41	41	10	41
7:	41	41	41	31	41	41	41	10

Core Affinity Example

```
#define _GNU_SOURCE
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#include <unistd.h>
```

```
#include <string.h>
```

```
#include <sched.h>
```

```
#include <mpi.h>
```

```
#include <omp.h>
```

```
static char *cpuset_to_cstr(cpu_set_t *mask, char *str)
{
    char *ptr = str;
    int i, j, entry_made = 0;
    for (i = 0; i < CPU_SETSIZE; i++) {
        if (CPU_ISSET(i, mask)) {
            int run = 0;
            entry_made = 1;
            for (j = i + 1; j < CPU_SETSIZE; j++) {
                if (CPU_ISSET(j, mask)) run++;
                else break;
            }
            if (!run)
                sprintf(ptr, "%d,", i);
            else if (run == 1) {
                sprintf(ptr, "%d,%d,", i, i + 1); i++;
            } else {
                sprintf(ptr, "%d-%d,", i, i + run);
                i += run;
            }
            while (*ptr != 0) ptr++;
        }
    }
    ptr -= entry_made;
    *ptr = 0;
    return(str);
}
```

```

int main(int argc, char *argv[])
{
    int thread_level, rank, thread, duration=0;
    cpu_set_t coremask;
    char clbuf[7 * CPU_SETSIZE], hnbuf[64];

    if (argc == 2) duration = atoi(argv[1]);
    /* MPI_Init(&argc, &argv); */
    MPI_Init_thread(&argc, &argv, MPI_THREAD_FUNNELED, &thread_level);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    memset(clbuf, 0, sizeof(clbuf));
    memset(hnbuf, 0, sizeof(hnbuf));
    gethostname(hnbuf, sizeof(hnbuf));
    #pragma omp parallel private(thread, coremask, clbuf)
    {
        thread = omp_get_thread_num();
        sched_getaffinity(0, sizeof(coremask), &coremask);
        cpuset_to_cstr(&coremask, clbuf);
        #pragma omp barrier
        printf("Hello from rank %d, thread %d, on %s. (core affinity = %s)\n",
              rank, thread, hnbuf, clbuf);
        sleep(duration);
    }
    MPI_Finalize();
    return(0);
}

```


Sample Batch Script for Hybrid Jobs on Edison

Batch script *hybrid.slurm*:

```
#!/bin/bash -l

#SBATCH -J hybrid
#SBATCH -p regular
#SBATCH -N 2
#SBATCH -t 4:00:00
#SBATCH -L SCRATCH
#SBATCH --mail-user=shaw@ucsc.edu
#SBATCH --mail-type=ALL

export OMP_NUM_THREADS=12
srun -n 4 -c 24 ./hybrid_hello.x
```

To submit the job:

```
sbatch hybrid.slurm
```

Comments:

```
### your favorite shell

### job name
### job queue
### request 2 node (48 physical cores)
### and 4 hours walltime
### request a license for scratch file system
### ask SLURM to send emails
### when jobs aborts, starts and ends

### 12 OpenMP threads per MPI task
### -n 4 : 4 MPI tasks (2 tasks per node)
### -c 24 : reserving 24 cores (12 physical plus
###         12 logical) per task;
###         although we only use physical cores
```

<http://www.nersc.gov/users/computational-systems/edison/running-jobs/example-batch-scripts/>

Sample Batch Script for Hybrid Jobs on Cori Phase I

Batch script *hybrid.slurm*:

```
#!/bin/bash -l

#SBATCH -J hybrid
#SBATCH -p regular
#SBATCH -C haswell
#SBATCH -N 2
#SBATCH -t 4:00:00
#SBATCH -L SCRATCH
#SBATCH --mail-user=shaw@ucsc.edu
#SBATCH --mail-type=ALL

export OMP_NUM_THREADS=16
srun -n 4 -c 32 ./hybrid_hello.x
```

To submit the job:

```
sbatch hybrid.slurm
```

Comments:

```
### your favorite shell

### job name
### job queue
### request a Haswell node
### request 2 node (64 physical cores)
### and 4 hours walltime
### request a license for scratch file system
### ask SLURM to send emails
### when jobs aborts, starts and ends

### 16 OpenMP threads per MPI task
### -n 4 : 4 MPI tasks (2 tasks per node)
### -c 32 : reserving 32 cores (16 physical plus
###          16 logical) per task;
###          although we only use physical cores
```

<http://www.nersc.gov/users/computational-systems/cori/running-jobs/example-batch-scripts/>

Sample Batch Script for Hybrid Jobs on Cori Phase III: snc4 + flat

```
#!/bin/bash -l
#SBATCH -N 2
#SBATCH -p regular
#SBATCH -C knl,snc4,flat
#SBATCH -S 4 # use core specialization
#SBATCH -t 3:00:00
#SBATCH -L SCRATCH,project

export OMP_NUM_THREADS=16 # 16 OpenMP threads per MPI task, using 64 cores in total
export OMP_PROC_BIND=true #"spread" is also good for Intel and CCE compilers
export OMP_PLACES=threads

# Add the following "sbcast" line here for jobs larger than 1500 MPI tasks:
# sbcast ./mycode.exe /tmp/mycode.exe

# 4 MPI ranks per node for a total of 8 MPI tasks
# executable linked with libmemkind
srun -n 8 -c 68 --cpu_bind=cores ./mycode.exe
```

<http://www.nersc.gov/users/computational-systems/cori/running-jobs/example-batch-scripts-for-knl/>

Sample Batch Script for Hybrid Jobs on Hyades

Batch script *hybrid.pbs*:

```
#!/bin/bash

#PBS -N hybrid
#PBS -q normal
#PBS -l nodes=2:ppn=16
#PBS -l walltime=0:10:00

cd $PBS_O_WORKDIR
cat $PBS_NODEFILE | sort | uniq > hosts.$PBS_JOBID

export OMP_NUM_THREADS=8
export I_MPI_PIN_DOMAIN=omp
export KMP_AFFINITY=compact

mpirun -machine hosts.$PBS_JOBID \
       -genv I_MPI_FABRICS shm:ofa \
       -n 4 -ppn 2 ./hybrid_hello.x
```

Comments:

```
### your favorite shell

### job name
### job queue
### request 2 node (16 cores per node)
### and 10 minutes walltime

### go to the work directory
### create a machine file

### 8 OpenMP threads per MPI task
### select process pinning scheme
### pin OpenMP threads

### run your hybrid executable
### -n 4: 4 MPI tasks in total
### -ppn 2: 2 MPI tasks per node
```

<https://software.intel.com/en-us/articles/hybrid-applications-intelmpi-openmp>

Try it out on Hyades

```
$ mpiicc -mt_mpi hybrid_hello.c -o hybrid_hello.x
```

```
$ qsub hybrid.pbs
```

```
$ cat hybrid.o125827
```

Hello from rank 0, thread 0, on astro-5-11.local. (core affinity = 0)

Hello from rank 0, thread 1, on astro-5-11.local. (core affinity = 2)

...

Hello from rank 1, thread 0, on astro-5-14.local. (core affinity = 0)

Hello from rank 1, thread 1, on astro-5-14.local. (core affinity = 2)

...

Hello from rank 2, thread 0, on astro-5-11.local. (core affinity = 1)


Hello from rank 2, thread 1, on astro-5-11.local. (core affinity = 3)

...

Hello from rank 3, thread 0, on astro-5-14.local. (core affinity = 1)

Hello from rank 3, thread 1, on astro-5-14.local. (core affinity = 3)

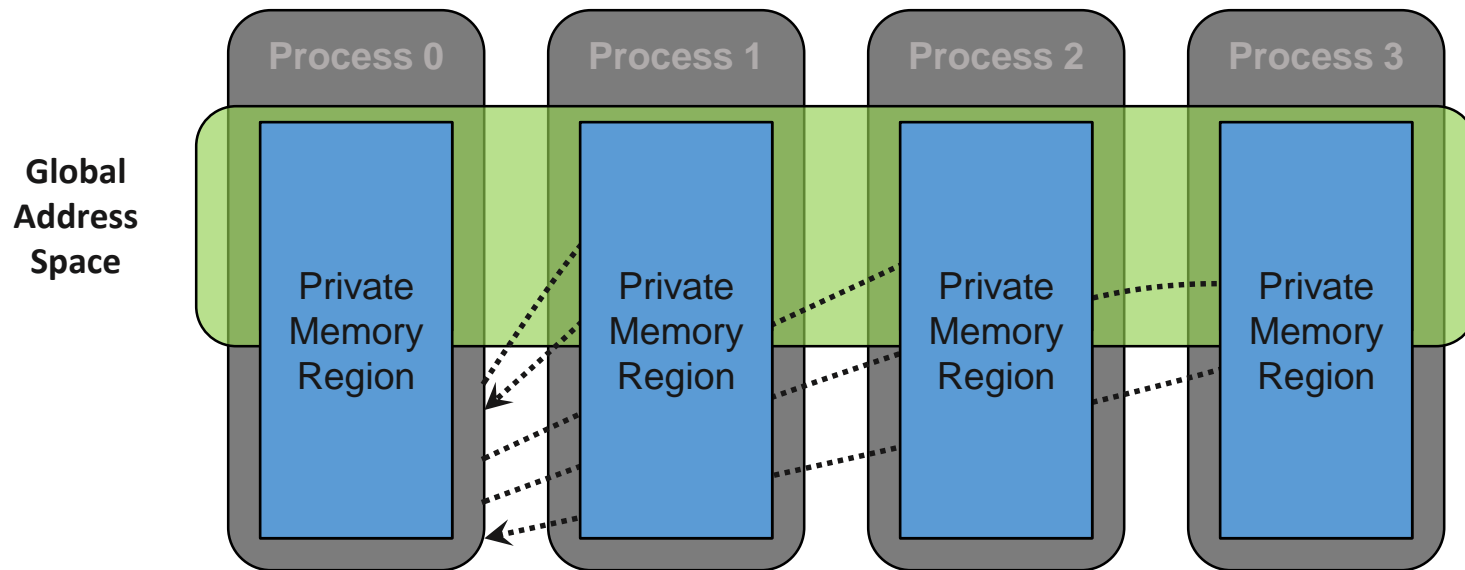
...



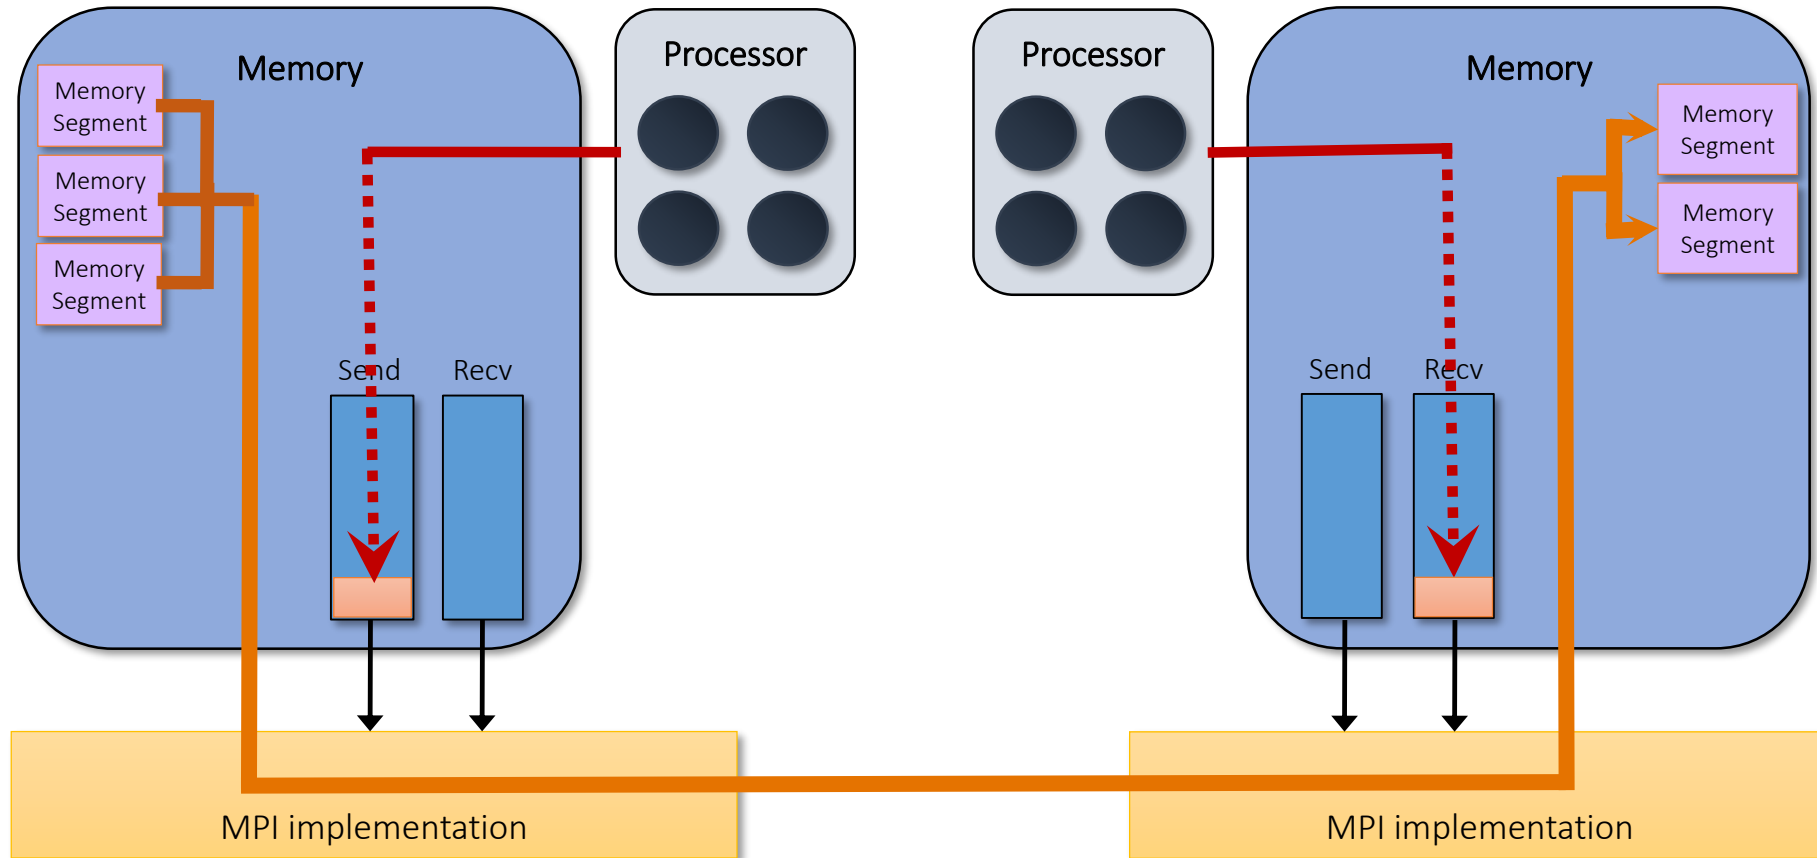
One-Sided Communications

The basic idea of one-sided communication models is to decouple data movement with process synchronization

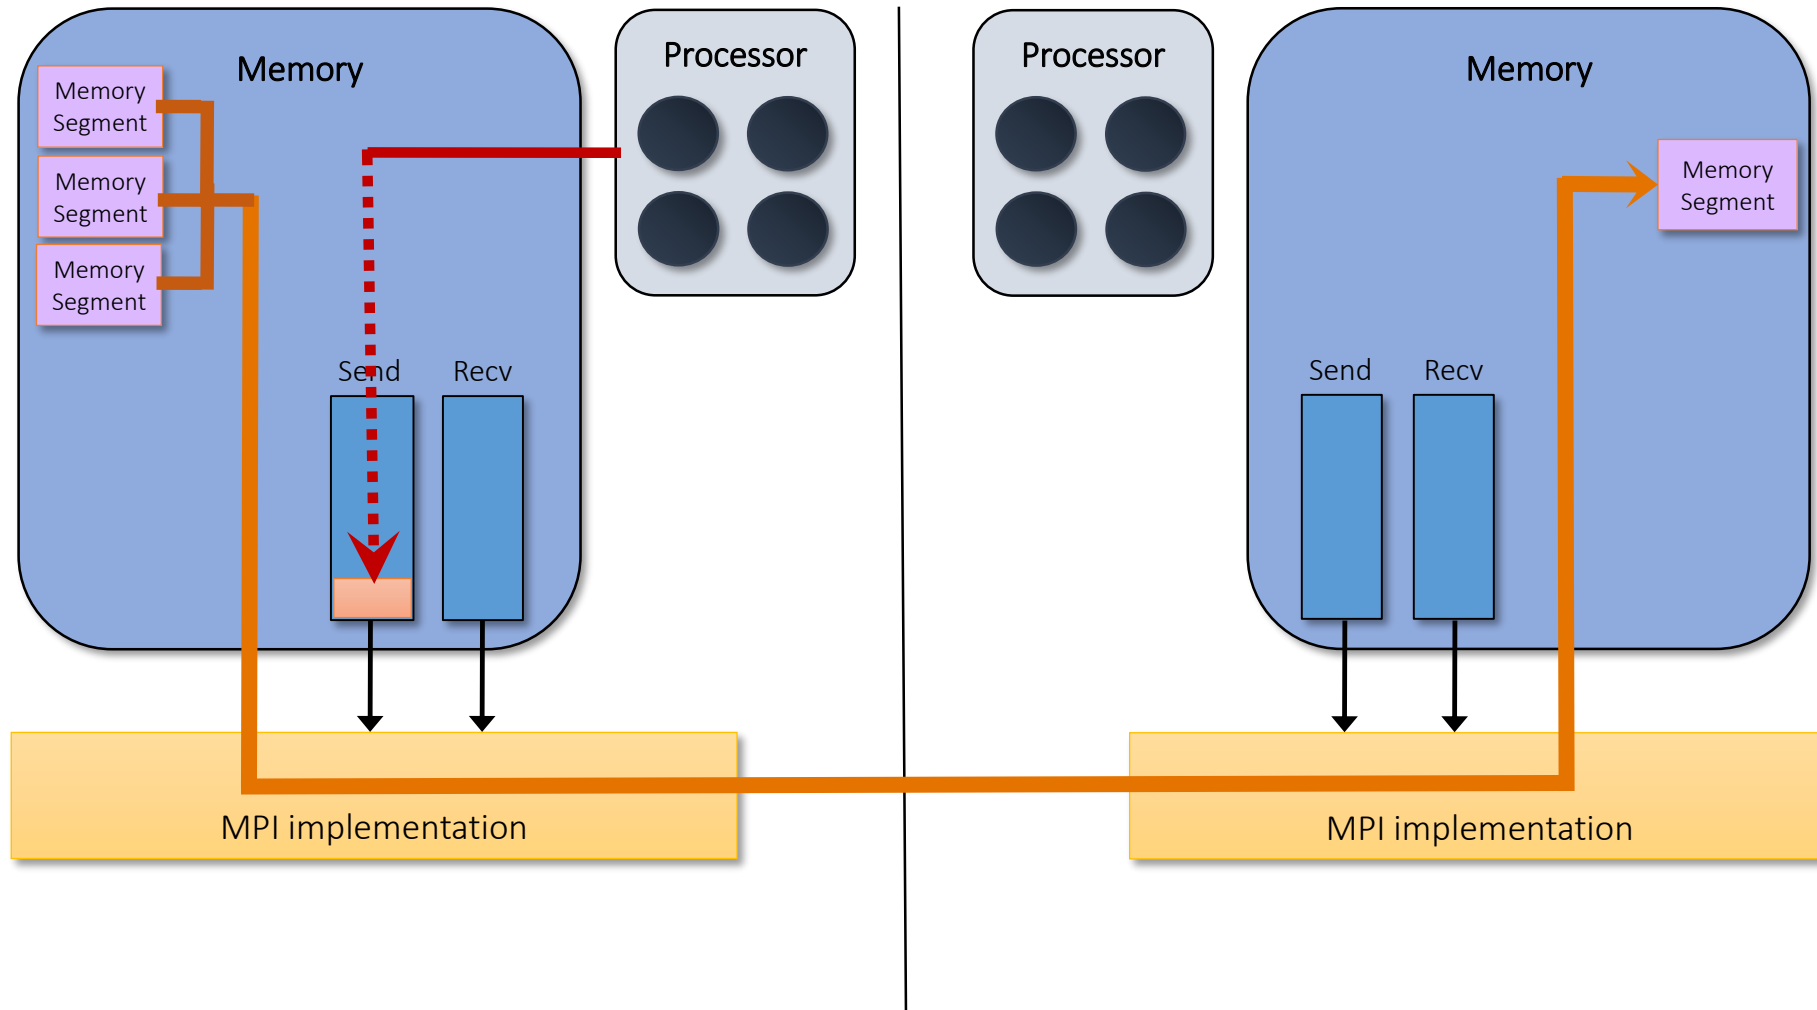
- Should be able move data without requiring that the remote process synchronize
- Each process exposes a part of its memory to other processes
- Other processes can *directly* read from or write to this memory



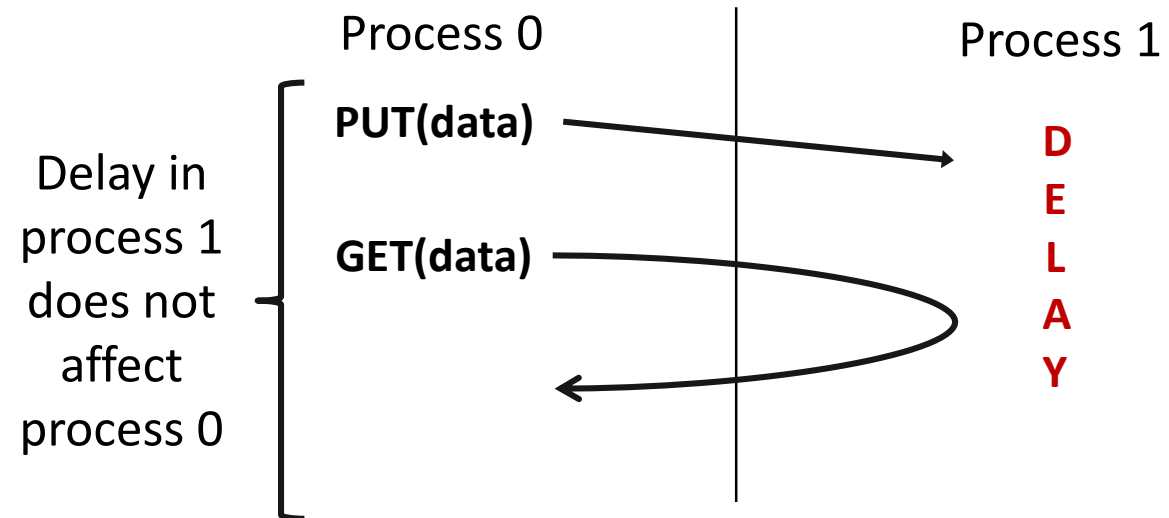
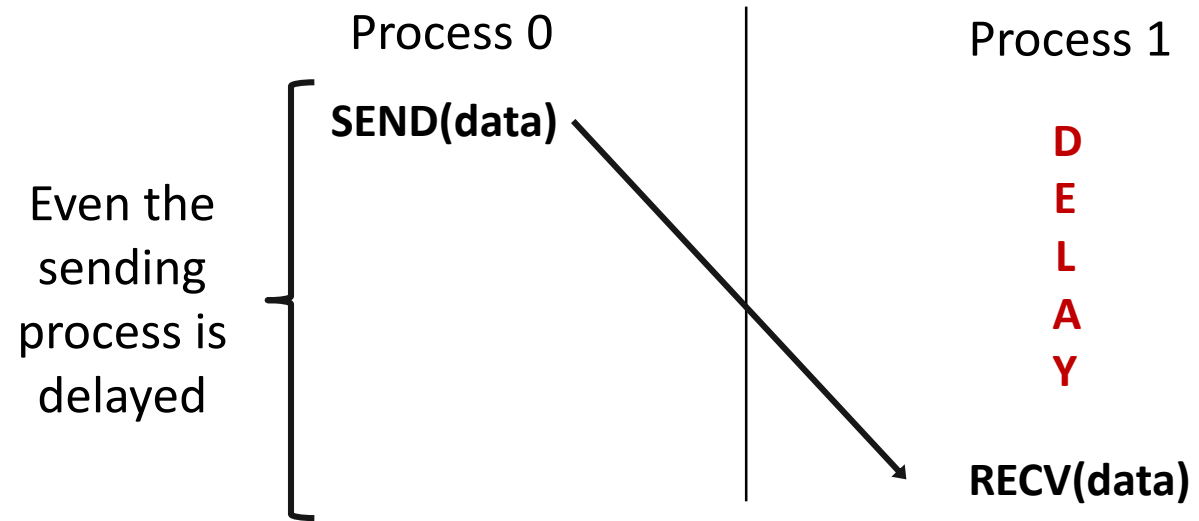
Two-Sided Communication Example



One-Sided Communication Example



One-Sided vs. Two-Sided



One-Sided Communications with MPI

- MPI-2 introduces routines for **one-sided communications**, aka **RMA** (remote memory access)
- MPI-3 adds new one-sided communication operations, to better handle different memory models
- Advantages of RMA operations:
 - Can do multiple data transfers with a single synchronization operation
 - Bypass tag matching – effectively precomputed as part of remote offset
 - Some irregular communication patterns can be more economically expressed
 - Can be significantly faster than send/receive on systems with hardware support for remote memory access, such as shared memory systems and RDMA

One-Sided Communication Terms

- **Origin process:** Process with the source buffer, initiates the operation
- **Target process:** Process with the destination buffer, does not explicitly call communication functions
- **Epoch:** Virtual time where operations are in flight. Data is consistent after new epoch is started.
 - Access epoch: rank acts as origin for RMA calls
 - Exposure epoch: rank acts as target for RMA calls
- **Ordering:** only for accumulate operations: order of messages between two processes (default: in order, can be relaxed)
- **Assert:** assertions about how One-Sided functions are used, “fast” optimization hints, cf. Info objects (slower)

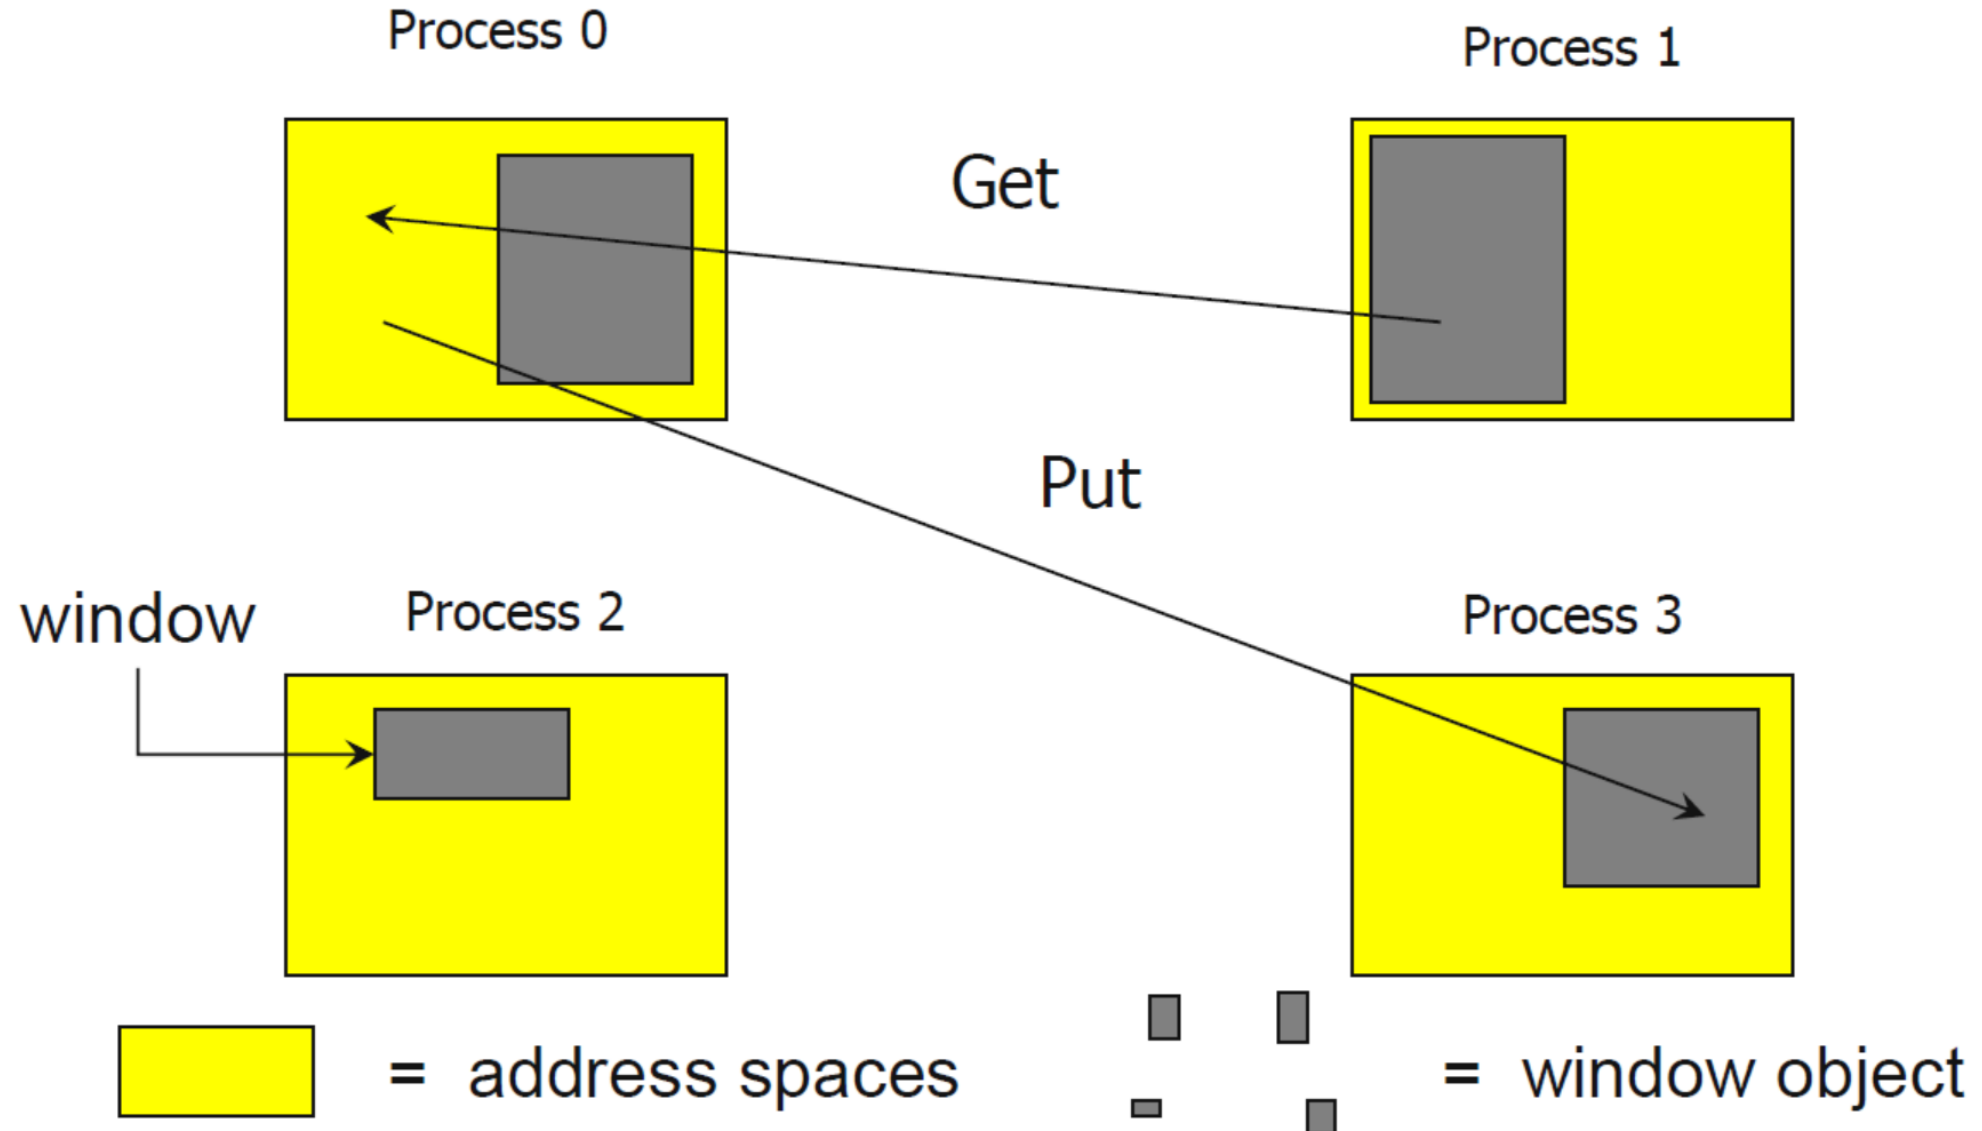
One-Sided Communication Overview

- Creation
 - Expose memory collectively – **Win_create**
 - Allocate exposed memory – **Win_allocate**
 - Dynamic memory exposure – **Win_create_dynamic**
- Communication
 - Data movement (**put, get, rput, rget**)
 - Accumulate (**acc, racc, get_acc, rget_acc, fetch&op, cas**)
- Synchronization
 - Active - Collective (**fence**); Group (**PSCW**)
 - Passive - P2P (**lock/unlock**); One epoch (**lock_all**)

Creating Public Memory

- Any memory created by a process is, by default, only locally accessible
`x = malloc(100);`
- Once the memory is created, the user has to make an *explicit* MPI call to declare a memory region as remotely accessible
 - MPI terminology for remotely accessible memory is a “**window**”
 - A group of processes *collectively* create a “**window**”
- Once a memory region is declared as remotely accessible, all processes in the window can read/write data to this memory without explicitly synchronizing with the target process

RMA Windows and Window Objects



Window Creation Models

Four window creation models:

- `MPI_Win_create`
 - You already have an allocated buffer that you would like to make remotely accessible
- `MPI_Win_allocate`
 - You want to create a buffer and directly make it remotely accessible
- `MPI_Win_create_dynamic`
 - You don't have a buffer yet, but will have one in the future
- `MPI_Win_allocate_shared`
 - You want multiple processes on the same node share a buffer

MPI_Win_create

C	<pre>int MPI_Win_create(void *base, MPI_Aint size, int disp_unit, MPI_Info info, MPI_Comm comm, MPI_Win *win)</pre>
Fortran	<pre>MPI_WIN_CREATE(BASE, SIZE, DISP_UNIT, INFO, COMM, WIN, IERROR) <type> BASE(*) INTEGER(KIND=MPI_ADDRESS_KIND) SIZE INTEGER DISP_UNIT, INFO, COMM, WIN, IERROR</pre>

- Expose a region of memory in an RMA window
- Collective call
- Arguments:
 - base – pointer to local data to expose
 - size – size of local data in bytes
 - disp_unit – local unit size for displacements, in bytes
 - info – info argument (handle)
 - comm – communicator (handle)
 - win – window object returned by the call (handle)

MPI_Win_free

C	<code>int MPI_Win_free(MPI_Win *win)</code>
Fortran	<code>MPI_WIN_FREE(WIN, IERROR)</code> <code>INTEGER WIN, IERROR</code>

- Frees the window object and returns a null handle
- Collective call

MPI_Win_create Example

```
#include "mpi.h"
int main(int argc, char ** argv)
{
    int *a;
    MPI_Win win;
    MPI_Init(&argc, &argv);
    /* create private memory */
    a = (void *) malloc(1000 * sizeof(int));
    /* use private memory like you normally would */
    a[0] = 1;  a[1] = 2;
    /* collectively declare memory as remotely accessible */
    MPI_Win_create(a, 1000*sizeof(int), sizeof(int),
                   MPI_INFO_NULL, MPI_COMM_WORLD, &win);
    /* Array 'a' is now accessibly by all processes in MPI_COMM_WORLD */
    MPI_Win_free(&win);
    free(a);
    MPI_Finalize();
    return 0;
}
```

MPI_Win_allocate

C	<code>int MPI_Win_allocate(MPI_Aint size, int disp_unit, MPI_Info info, MPI_Comm comm, void *baseptr, MPI_Win *win)</code>
Fortran	<code>MPI_WIN_ALLOCATE(IZE, DISP_UNIT, INFO, COMM, BASEPTR, WIN, IERROR) INTEGER(KIND=MPI_ADDRESS_KIND) SIZE, BASEPTR INTEGER DISP_UNIT, INFO, COMM, WIN, IERROR</code>

- Allocates memory and returns a window object for RMA operations
- Collective call
- Arguments:
 - size – size of window in bytes
 - disp_unit – local unit size for displacements, in bytes
 - info – info argument (handle)
 - comm – communicator (handle)
 - baseptr – initial address of window
 - win – window object returned by the call (handle)

MPI_Win_allocate Example

```
#include "mpi.h"
int main(int argc, char ** argv)
{
    int *a;
    MPI_Win win;
    MPI_Init(&argc, &argv);

    /* collectively create remotely accessible memory in the window */
    MPI_Win_allocate(1000, sizeof(int), MPI_INFO_NULL,
                    MPI_COMM_WORLD, &a, &win);

    /* Array 'a' is now accessibly by all processes in MPI_COMM_WORLD */
    MPI_Win_free(&win);

    MPI_Finalize();
    return 0;
}
```

MPI_Win_create_dynamic

C	<pre>int MPI_Win_create_dynamic(MPI_Info info, MPI_Comm comm, MPI_Win *win)</pre>
Fortran	<pre>MPI_WIN_CREATE_DYNAMIC(INFO, COMM, WIN, IERROR) INTEGER INFO, COMM, WIN, IERROR</pre>

- Creates an RMA window, to which data can later be attached
- Collective call
- Application can dynamically attach memory to this window
- Application can access data on this window only after a memory region has been attached

MPI_Win_create_dynamic Example

```
#include "mpi.h"
int main(int argc, char ** argv)
{
    int *a;
    MPI_Win win;
    MPI_Init(&argc, &argv);
    MPI_Win_create_dynamic(MPI_INFO_NULL, MPI_COMM_WORLD, &win);
    /* create private memory */
    a = (void *) malloc(1000 * sizeof(int));
    /* use private memory like you normally would */
    a[0] = 1;  a[1] = 2;
    /* locally declare memory as remotely accessible */
    MPI_Win_attach(win, a, 1000);
    /* Array 'a' is now accessible from all processes in MPI_COMM_WORLD */
    /* undeclare public memory */
    MPI_Win_detach(win, a);
    MPI_Win_free(&win);
    MPI_Finalize(); return 0;
}
```

Data Movement

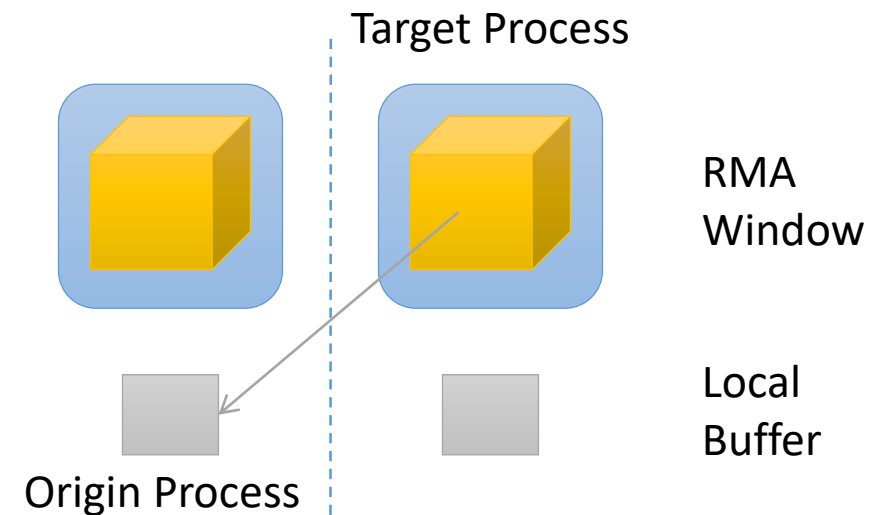
MPI provides ability to read, write and atomically modify data in remotely accessible memory regions

- MPI_Get
- MPI_Put
- MPI_Accumulate
- MPI_Get_accumulate
- MPI_Compare_and_swap
- MPI_Fetch_and_op

MPI_Get

C	<pre>int MPI_Get(void *origin_addr, int origin_count, MPI_Datatype origin_datatype, int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype, MPI_Win win)</pre>
Fortran	<pre>MPI_GET(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR) <type> ORIGIN_ADDR(*) INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR</pre>

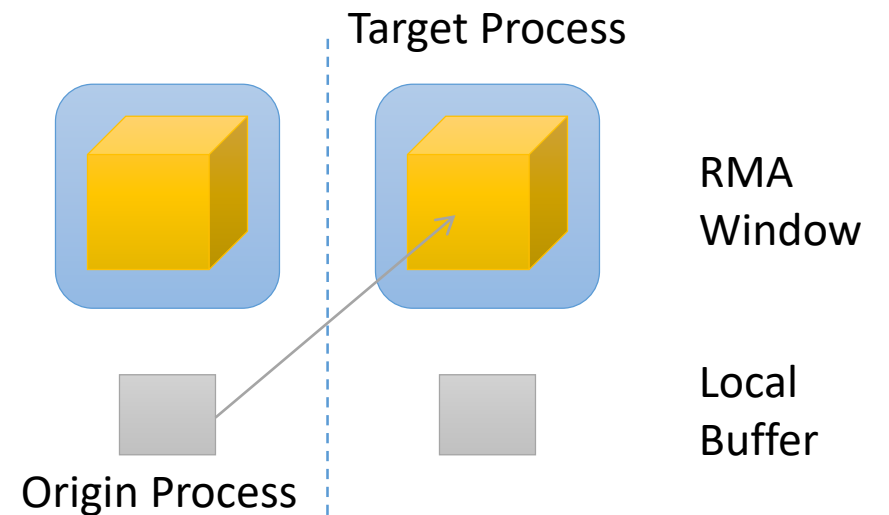
- Copies data from the **target** memory to the **origin**
- Separate data description triples for **origin** and **target**



MPI_Put

C	<pre>int MPI_Put(const void *origin_addr, int origin_count, MPI_Datatype origin_datatype, int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype, MPI_Win win)</pre>
Fortran	<pre>MPI_PUT(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR) <type> ORIGIN_ADDR(*) INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT, TARGET_DATATYPE, WIN, IERROR</pre>

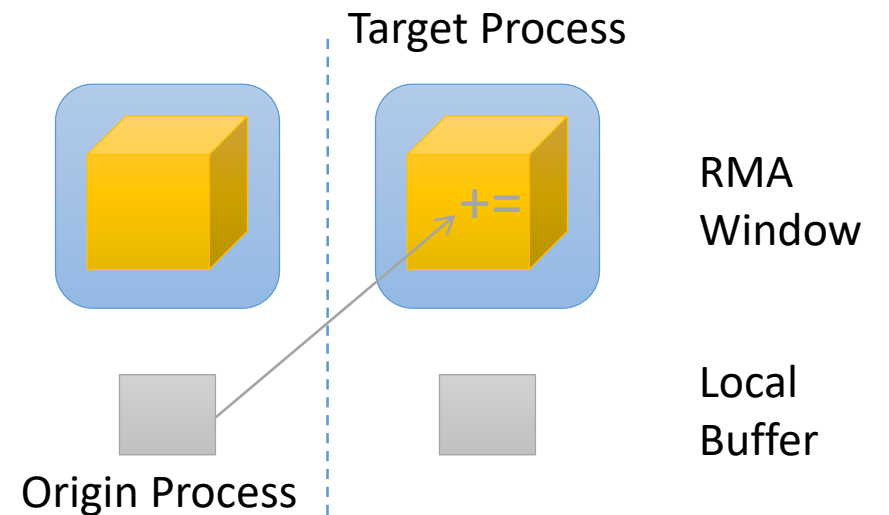
- Copies data from the **origin** memory to the **target**
- Separate data description triples for **origin** and **target**
- Same arguments as MPI_Get



MPI_Accumulate

C	<pre>int MPI_Accumulate(const void *origin_addr, int origin_count, MPI_Datatype origin_datatype, int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype, MPI_Op op, MPI_Win win)</pre>
Fortran	<pre>MPI_ACCUMULATE(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR) <type> ORIGIN_ADDR(*) INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_RANK, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR</pre>


- Combines the contents of the **origin** buffer with that of a **target** buffer
- Any of the predefined operations for **MPI_Reduce** can be used. User-defined functions cannot be used.
- Atomic *put* with op = **MPI_REPLACE**



MPI_Get_accumulate

C	<pre>int MPI_Get_accumulate(const void *origin_addr, int origin_count, MPI_Datatype origin_datatype, void *result_addr, int result_count, MPI_Datatype result_datatype, int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype, MPI_Op op, MPI_Win win)</pre>
Fortran	<pre>MPI_GET_ACCUMULATE(ORIGIN_ADDR, ORIGIN_COUNT, ORIGIN_DATATYPE, RESULT_ADDR, RESULT_COUNT, RESULT_DATATYPE, TARGET_RANK, TARGET_DISP, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR) <type> ORIGIN_ADDR, RESULT_ADDR(*) INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP INTEGER ORIGIN_COUNT, ORIGIN_DATATYPE, TARGET_COUNT, TARGET_DATATYPE, TARGET_RANK, TARGET_COUNT, TARGET_DATATYPE, OP, WIN, IERROR</pre>

- Combines the contents of the **origin** buffer with that of a **target** buffer and returns the **target** buffer value
- Predefined ops only, no user-defined!
- Atomic *get* with op = **MPI_NO_OP**
- Atomic *swap* with op = **MPI_REPLACE**



MPI_Compare_and_swap

C	<pre>int MPI_Compare_and_swap(const void *origin_addr, const void *compar_addr, void *result_addr, MPI_Datatype datatype, int target_rank, MPI_Aint target_disp, MPI_Win win)</pre>
Fortran	<pre>MPI_COMPARE_AND_SWAP(ORIGIN_ADDR, COMPARE_ADDR, RESULT_ADDR, DATATYPE, TARGET_RANK, TARGET_DISP, WIN, IERROR) <type> ORIGIN_ADDR, COMPARE_ADDR, RESULT_ADDR(*) INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP INTEGER DATATYPE, TARGET_RANK, WIN, IERROR</pre>

- Performs RMA compare-and-swap
- This function compares one element of type datatype in the compare buffer *compare_addr* with the buffer at offset *target_disp* in the target window specified by *target_rank* and *win* and replaces the value at the target with the value in the origin buffer *origin_addr* if the compare buffer and the target buffer are identical. The original value at the target is returned in the buffer *result_addr*.
- Atomic swap if target value is equal to compare value

MPI_Fetch_and_op

C	<pre>int MPI_Fetch_and_op(const void *origin_addr, void *result_addr, MPI_Datatype datatype, int target_rank, MPI_Aint target_disp, MPI_Op op, MPI_Win win)</pre>
Fortran	<pre>MPI_FETCH_AND_OP(ORIGIN_ADDR, RESULT_ADDR, DATATYPE, TARGET_RANK, TARGET_DISP, OP, WIN, IERROR) <type> ORIGIN_ADDR, RESULT_ADDR(*) INTEGER(KIND=MPI_ADDRESS_KIND) TARGET_DISP INTEGER DATATYPE, TARGET_RANK, OP, WIN, IERROR</pre>

- Combines the contents of the **origin** buffer with that of a **target** buffer and returns the **target** buffer value
- Simpler version of **MPI_Get_accumulate**
 - All buffers share a single predefined datatype
 - No count argument (it's always 1)
 - Simpler interface allows hardware optimization

Ordering of RMA Operations

- No guaranteed ordering for Put/Get operations
- Result of concurrent Puts to the same location undefined
- Result of Get and concurrent Put/Accumulate undefined
 - Can be garbage in both cases
- Result of concurrent accumulate operations to the same location are defined according to the order in which they occurred
 - Atomic put: **Accumulate** with op = **MPI_REPLACE**
 - Atomic get: **Get_accumulate** with op = **MPI_NO_OP**
- Accumulate operations from a given process are ordered by default
 - User can tell the MPI implementation that ordering is not required as optimization hint
 - You can ask for only the needed orderings, e.g., RAW (read-after-write), WAR, RAR, or WAW

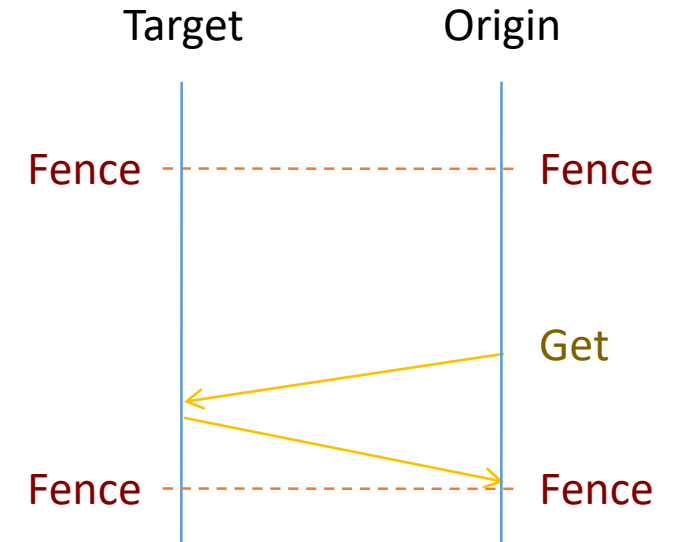
RMA Synchronization Models

- RMA data visibility
 - When is a process allowed to read/write remotely accessible memory?
 - When is data written by process X is available for process Y to read?
 - RMA synchronization models define these semantics
- Three synchronization models provided by MPI:
 - **Fence** (*active* target)
 - **Post-Start-Complete-Wait** (generalized *active* target)
 - **Lock/Unlock** (*passive* target)
- Data accesses occur within “epochs”
 - Access epochs: contain a set of operations issued by an origin process
 - Exposure epochs: enable remote processes to access and/or update a target’s window
 - Epochs define ordering and completion semantics
 - Synchronization models provide mechanisms for establishing epochs
 - E.g., starting, ending, and synchronizing epochs

Fence: Active Target Synchronization

C	<code>int MPI_win_fence(int assert, MPI_win win)</code>
Fortran	<code>MPI_WIN_FENCE(ASSERT, WIN, IERROR)</code> <code>INTEGER ASSERT, WIN, IERROR</code>

- *Collective* synchronization model
- Starts and ends access and exposure epochs on all processes in the window
- All processes in group of “win” do an **MPI_Win_fence** to open an epoch
- Everyone can issue Put/Get operations to read/write data
- Everyone does an **MPI_Win_fence** to close the epoch
- All operations complete at the second fence synchronization



Example: Calculating π with MPI RMA

```
#include "mpi.h"
#include <math.h>
int main(int argc, char *argv[]) {
    int n, myid, numprocs, i;
    double PI25DT = 3.141592653589793238462643;
    double mypi, pi, h, sum, x;
    MPI_WIN nwin, piwin;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    if (myid == 0) {
        MPI_Win_create(&n, sizeof(int), 1, MPI_INFO_NULL,
                      MPI_COMM_WORLD, &nwin);
        MPI_Win_create(&pi, sizeof(double), 1, MPI_INFO_NULL,
                      MPI_COMM_WORLD, &piwin);
    } else {
        MPI_Win_create(MPI_BOTTOM, 0, 1, MPI_INFO_NULL,
                      MPI_COMM_WORLD, &nwin);
        MPI_Win_create(MPI_BOTTOM, 0, 1, MPI_INFO_NULL,
                      MPI_COMM_WORLD, &piwin);
    }
}
```


$$\frac{\pi}{4} = \tan^{-1}(-1) = \int_0^1 \frac{dx}{1+x^2}$$

MPI_BOTTOM: a predefined constant, indicating the bottom of the address space

```

while (1) {
    if (myid == 0) {
        printf("Enter the number of intervals: (0 quits) ");
        scanf("%d",&n);
        pi = 0.0;
    }
    MPI_Win_fence(0, nwin);
    if (myid != 0) MPI_Get(&n, 1, MPI_INT, 0, 0, 1, MPI_INT, nwin);
    MPI_Win_fence(0, nwin);
    if (n == 0)
        break;
    else {
        h = 1.0 / (double) n;
        sum = 0.0;
        for (i = myid + 1; i <= n; i += numprocs) {
            x = h * ((double)i - 0.5);
            sum += (4.0 / (1.0 + x*x));
        }
        mypi = h * sum;
        MPI_Win_fence(0, piwin);
        MPI_Accumulate(&mypi, 1, MPI_DOUBLE, 0, 0, 1, MPI_DOUBLE, MPI_SUM, piwin);
        MPI_Win_fence(0, piwin);
        if (myid == 0)
            printf("pi is approximately %.16f, Error is %.16f\n", pi, fabs(pi - PI25DT));
    }
}
MPI_Win_free(&nwin); MPI_Win_free(&piwin);
MPI_Finalize();
return 0;
}

```

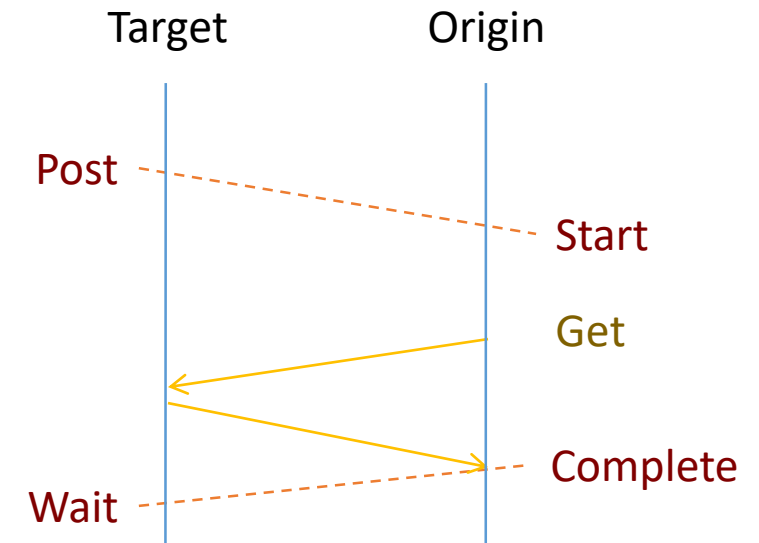


$$\frac{\pi}{4} = \tan^{-1}(-1) = \int_0^1 \frac{dx}{1+x^2}$$


PSCW: Generalized Active Target Synchronization

```
int MPI_win_post/start(MPI_Group group, int assert, MPI_Win win)
int MPI_win_complete/wait(MPI_Win win)
```

- Like **Fence**, but origin and target specify who they communicate with
- Target: Exposure epoch
 - Opened with **MPI_Win_post**
 - Closed by **MPI_Win_wait**
- Origin: Access epoch
 - Opened by **MPI_Win_start**
 - Closed by **MPI_Win_complete**
- All synchronization operations may block, to enforce P-S/C-W ordering
 - Processes can be both origins and targets



Passive Target Mode



- Passive mode: One-sided, *asynchronous* communication
 - Target does not participate in communication operation
- Shared memory like model

Passive Target Synchronization

```
int MPI_Win_lock(int lock_type, int rank, int assert, MPI_Win win)
int MPI_Win_unlock(int rank, MPI_Win win)
int MPI_Win_flush/flush_local(int rank, MPI_Win win)
```

- **Lock/Unlock:** Begin/end passive mode epoch
 - Target process does not make a corresponding MPI call
 - Can initiate multiple passive target epochs to different processes
 - Concurrent epochs to same process not allowed (affects threads)
- **Lock type**
 - **MPI_LOCK_SHARED:** Other processes using shared can access concurrently
 - **MPI_LOCK_EXCLUSIVE:** No other processes can access concurrently
- **Flush:** Remotely complete RMA operations to the target process
 - After completion, data can be read by target process or a different process
- **Flush_local:** Locally complete RMA operations to the target process

Lock is not Lock

- The name “Lock” is unfortunate
 - **Lock** is really “begin epoch”
 - **Unlock** is really “end epoch”
- An MPI “Lock” does not establish a critical section or mutual exclusion
 - With **MPI_LOCK_EXCLUSIVE** the RMA operations have exclusive access to the data they access/update during the time that they access the remote window
- This is very different than a “lock” in the sense of a thread lock

MPI_Win_lock Example

```
#include "mpi.h"
#include "stdio.h"

/* tests passive target RMA on 2 processes */

#define SIZE1 100
#define SIZE2 200

int main(int argc, char *argv[])
{
    int rank, nprocs, A[SIZE2], B[SIZE2], i;
    MPI_Win win;

    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if (nprocs != 2) {
        printf("Run this program with 2 processes\n")
        MPI_Abort(MPI_COMM_WORLD,1);
    }
}
```

```

if (rank == 0) {
    for (i=0; i<SIZE2; i++) A[i] = B[i] = i;
    MPI_win_create(NULL, 0, 1, MPI_INFO_NULL, MPI_COMM_WORLD, &win);

    for (i=0; i<SIZE1; i++) {
        MPI_win_lock(MPI_LOCK_SHARED, 1, 0, win);
        MPI_Put(A+i, 1, MPI_INT, 1, i, 1, MPI_INT, win);
        MPI_win_unlock(1, win);
    }

    for (i=0; i<SIZE1; i++) {
        MPI_win_lock(MPI_LOCK_SHARED, 1, 0, win);
        MPI_Get(B+i, 1, MPI_INT, 1, SIZE1+i, 1, MPI_INT, win);
        MPI_win_unlock(1, win);
    }

    MPI_win_free(&win);

else { /* rank=1 */
    for (i=0; i<SIZE2; i++) B[i] = (-4)*i;
    MPI_win_create(B, SIZE2*sizeof(int), sizeof(int), MPI_INFO_NULL, MPI_COMM_WORLD, &win);
    MPI_win_free(&win);
}
MPI_Finalize();
return 0;
}

```


Advanced Passive Target Synchronization

```
int MPI_win_lock_all(int assert, MPI_Win win)
int MPI_win_unlock_all(MPI_Win win)
int MPI_win_flush_all/flush_local_all(MPI_Win win)
```

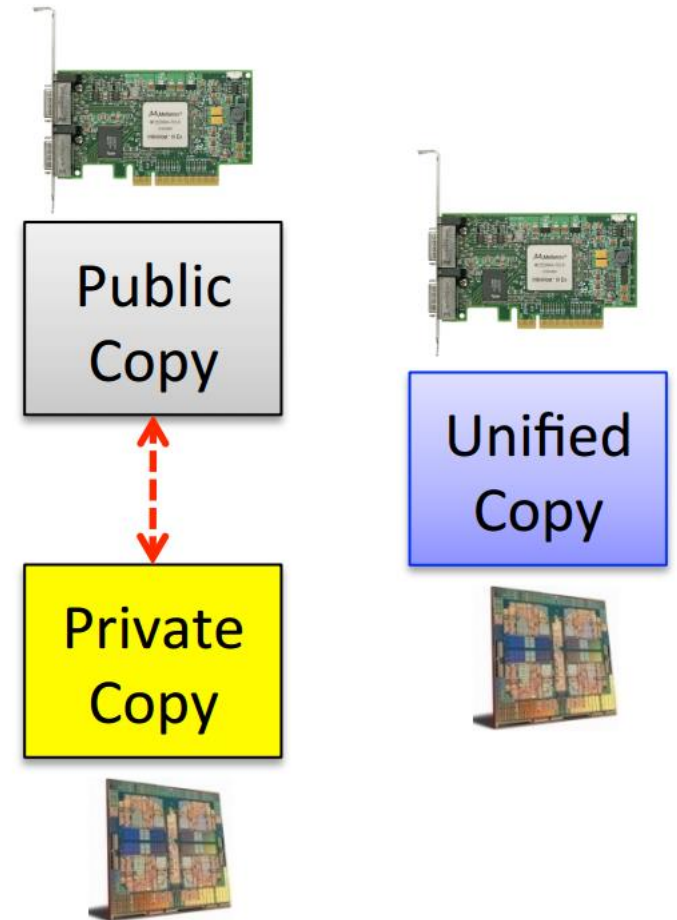
- **Lock_all**: starts an RMA access epoch locking access to all processes in the window, with a lock type of **MPI_LOCK_SHARED**
 - Expected usage is long-lived: lock_all, put/get, flush, ..., unlock_all
- **Flush_all**: remotely completes RMA operations to all processes
- **Flush_local_all**: locally completes RMA operations to all processes

Which Synchronization Mode to Use?

- RMA communication has low overheads versus send/recv
 - Two-sided: Matching, queuing, buffering, unexpected receives, etc.
 - One-sided: No matching, no buffering, always ready to receive
 - Utilize RDMA provided by high-speed interconnects (e.g., InfiniBand)
- Active mode: bulk synchronization
 - E.g., ghost cell exchange
- Passive mode: asynchronous data movement
 - Useful when dataset is large, requiring memory of multiple nodes
 - And when data access and synchronization pattern is dynamic
 - Common use case: distributed, shared arrays
- Passive target locking mode
 - Lock/unlock – Useful when exclusive epochs are needed
 - Lock_all/unlock_all – Useful when only shared epochs are needed

Memory Models

- MPI-3 provides two memory models: **separate** and **unified**
 - Type is attached as attribute to window
- MPI-2: **Separate** Model
 - Logical public and private copies
 - MPI provides software coherence between window copies
 - Extremely portable, to systems that don't provide hardware coherence
- MPI-3: New **Unified** Model
 - Single copy of the window
 - System must provide coherence
 - Superset of separate semantics
 - E.g., allows concurrent local/remote access
 - Provides access to full performance potential of hardware



Separate Semantics

	Load	Store	Get	Put	Acc
Load	OVL+NOVL	OVL+NOVL	OVL+NOVL	X	X
Store	OVL+NOVL	OVL+NOVL	NOVL	X	X
Get	OVL+NOVL	NOVL	OVL+NOVL	NOVL	NOVL
Put	X	X	NOVL	NOVL	NOVL
Acc	X	X	NOVL	NOVL	OVL+NOVL

This matrix shows the compatibility of MPI-RMA operations when two or more processes access a window at the same target concurrently.

OVL – Overlapping operations permitted

NOVL – Nonoverlapping operations permitted

X – Combining these operations is OK, but data might be garbage

Unified Semantics

	Load	Store	Get	Put	Acc
Load	OVL+NOVL	OVL+NOVL	OVL+NOVL	NOVL	NOVL
Store	OVL+NOVL	OVL+NOVL	NOVL	NOVL	NOVL
Get	OVL+NOVL	NOVL	OVL+NOVL	NOVL	NOVL
Put	NOVL	NOVL	NOVL	NOVL	NOVL
Acc	NOVL	NOVL	NOVL	NOVL	OVL+NOVL

This matrix shows the compatibility of MPI-RMA operations when two or more processes access a window at the same target concurrently.

OVL – Overlapping operations permitted

NOVL – Nonoverlapping operations permitted

Further Readings

- **Using MPI**, 3rd Edition, *by* William Gropp, Ewing Lusk & Anthony Skjellum, MIT Press, 2014
<http://ieeexplore.ieee.org/xpl/bkabstractplus.jsp?bkn=6981847>
- **Using Advanced MPI**, *by* William Gropp, Torsten Hoefler, Rajeev Thakur & Ewing Lusk, MIT Press, 2014
<http://ieeexplore.ieee.org/xpl/bkabstractplus.jsp?bkn=6981848>
- Using MPI and Using Advanced MPI:
<http://wgropp.cs.illinois.edu/usingmpiweb/>