

AMS 250: An Introduction to High Performance Computing

PGAS



Shawfeng Dong

shaw@ucsc.edu

(831) 502-7743

Applied Mathematics & Statistics
University of California, Santa Cruz

Outline

- PGAS Overview
- Coarray Fortran (CAF)
- Unified Parallel C (UPC)
- UPC++

Recap: Parallel Programming Models

- Process Interaction:
 - Shared Memory
 - Message Passing
 - **Partitioned Global Address Space (PGAS)**
- Problem Decomposition
 - Task Parallelism
 - Data Parallelism


PGAS (Partitioned Global Address Space)

- A **global** memory **address space** that is logically **partitioned** and a portion of it is local to each process or thread
- One-sided communication
- Explicit synchronization, as opposed to (mostly) implicit for MPI
- PGAS libraries:
 - MPI One-Sided (RMA), OpenSHMEM, Global Arrays, UPC++, etc.
- PGAS languages:
 - UPC (Unified Parallel C) – an extension to C
 - CAF (Coarray Fortran) – part of Fortran 2008 standard
- APGAS (asynchronous partitioned global address space) languages, which permit both local and remote asynchronous task creation:
 - Chapel: <http://chapel.cray.com/>
 - X10: <http://x10-lang.org/>

Recap: One-Sided Communications


The basic idea of one-sided communication models is to decouple data movement with process synchronization

- Should be able move data without requiring that the remote process synchronize
- Each process exposes a part of its memory to other processes
- Other processes can *directly* read from or write to this memory



Partitioned Global Address Space

- **Global Address Space:** thread/process may directly read/write remote data
 - Convenience of shared memory
- **Partitioned:** data is designated as local or global
 - Locality and scalability of message passing



Fortran 2008

- Fortran 2008 (latest draft)
<http://www.j3-fortran.org/doc/year/10/10-007.pdf>
- Fortran 2008 is a natively parallel language
 - SPMD programming model
 - Simple syntax for one-sided communication, using normal assignment statements
- Executable is replicated across processors (MPI-like)
- Each instance is called an “**IMAGE**”
- Each image has its own data objects
- Each image executes *asynchronously* except when syncs are indicated

“Hello, world!” in Coarray Fortran (CAF)

```
program caf_hello

  character*80 hostname

  call hostnm(hostname)

  write(*,*) "Hello from CAF image ", &
             this_image(), &
             "running on ", trim(hostname), &
             " out of ", num_images()

end program caf_hello
```


Compared to MPI “Hello, world!” in Fortran 90

```
program hello

  use mpi
  implicit none
  integer :: ierr, rank, size

  call MPI_INIT(ierr)
  call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
  call MPI_COMM_SIZE(MPI_COMM_WORLD, size, ierr)
  print *, "Hello, world! I am process ", rank, " of ", size
  call MPI_FINALIZE(ierr)

end program hello
```

Coarray Fortran on Hyades (Intel Compilers)

- Intel Fortran compiler supports parallel programming using coarrays as defined in the Fortran 2008 standard
- A CAF program can be built to run on:
 - either a *shared memory* system
 - or a *distributed memory* cluster
- To compile the example CAF program for a shared memory system:
`ifort -coarray caf_hello.f90 -o caf_hello.x`
- By default, when a CAF program is compiled with the Intel compiler, the invocation creates as many images as there are processor cores on the host platform. For example, on the master node:

```
$ ./caf_hello.x
Hello from CAF image 6 running on hyades.ucsc.edu out of 32
Hello from CAF image 7 running on hyades.ucsc.edu out of 32
Hello from CAF image 14 running on hyades.ucsc.edu out of 32
...
```

CAF on Shared Memory System

There are 2 ways to control the number of images on a shared memory system:

- Use the `coarray-num-images=N` compiler option. E.g.:

```
$ ifort -coarray -coarray-num-images=2 caf_hello.f90 -o caf_hello.x
```

```
$ ./caf_hello.x
```

```
Hello from image      1 running on hyades.ucsc.edu out of      2
```

```
Hello from image      2 running on hyades.ucsc.edu out of      2
```

- Use the environment variable `FOR_COARRAY_NUM_IMAGES`. E.g.:

```
$ export FOR_COARRAY_NUM_IMAGES=4
```

```
$ ./caf_hello.x
```

```
Hello from CAF image   1 running on hyades.ucsc.edu out of   4
```

```
Hello from CAF image   2 running on hyades.ucsc.edu out of   4
```

```
Hello from CAF image   3 running on hyades.ucsc.edu out of   4
```

```
Hello from CAF image   4 running on hyades.ucsc.edu out of   4
```

CAF on Distributed Memory Cluster

- Set up a machine file, e.g. (*hosts*):
gpu-1:2
gpu-2:2
- Set up a CAF configuration file, e.g. (*cafconfig.txt*):
-genvall -genv I_MPI_FABRICS shm:ofa -machinefile hosts -n 4 ./caf_hello.dist
- Compile the example CAF program for a distributed memory cluster:
ifort -coarray=distributed -coarray-config-file=cafconfig.txt caf_hello.f90 -o caf_hello.dist
- Run the CAF application:
\$ mpdboot -n 3 -f hosts
\$./caf_hello.dist
Hello from CAF image 1 running on gpu-1.local out of 4
Hello from CAF image 2 running on gpu-1.local out of 4
Hello from CAF image 3 running on gpu-2.local out of 4
Hello from CAF image 4 running on gpu-2.local out of 4
\$ mpdallexit

CAF on Cori

- CAF is supported on Cori through 2 different implementations: Cray CAF and Intel CAF

<http://www.nersc.gov/users/computational-systems/cori/programming/additional-programming-models/>

- Cray CAF

Switch to the Cray compiler environment:

```
cori09> module swap PrgEnv-intel PrgEnv-cray
```

Supply the '**-h caf**' option when calling ftn:

```
cori09> ftn -h caf caf_hello.f90 -o caf_hello.x
```

```
cori09> salloc -N 2 -t 10:00 -p debug -C haswell
```

```
nid00461> ulimit -v unlimited # may not be necessary
```

```
nid00461> srun -n 64 ./caf_hello.x
```

Coarrays in Fortran 2008

- The array syntax of Fortran is extended with additional trailing subscripts in square brackets (`[]`) to provide a concise representation of references to data that is spread across images:
 - e.g., `real :: a(3)[*]`
- Any time a coarray appears without `[]`, the reference is to the data on the local image
- The number inside the `[]` can reference any image in the job, including itself
- If a reference with `[]` appears to the right of the `=`, it is often called a “get”
 - e.g., `b(:) = a(:)[ri]`
- If a reference with `[]` appears to the left of the `=`, it is often called a “put”
 - e.g., `a(:)[ri] = b(:)`

Fortran 2008 Parallel Programming

- Declaration and allocation

```
real(8), ALLOCATABLE :: rcvbuf(:, :)[*]  
! Allocate m*n elements on each processor  
ALLOCATE( rcvbuf(m,n)[*] )
```

- Reference

```
! Put data from my local buf into rcvbuf on image k  
rcvbuf(:, :)[k] = localbuf(:, :)
```

- PE (processing elements) information

```
this_image(), num_images()
```

- Synchronization

```
sync all  
sync images(array_of_images)
```

Array Example

`real :: a(3)`

a(1)
a(2)
a(3)

Image 1

a(1)
a(2)
a(3)

Image 2

a(1)
a(2)
a(3)

Image 3

a(1)
a(2)
a(3)

Image 4

Coarray Example

```
real :: a(3)[*]
```

a(1)[1]
a(2)[1]
a(3)[1]

Image 1

a(1)[2]
a(2)[2]
a(3)[2]

Image 2

a(1)[3]
a(2)[3]
a(3)[3]

Image 3

a(1)[4]
a(2)[4]
a(3)[4]

Image 4

Fortran 2008 Synchronization

Explicit statements:

- sync all
- sync images (array_of_images)
- sync memory
- critical / end critical
- lock / unlock

Implicit synchronization:


- allocation of a coarray
- deallocation of a coarray (either explicit or implicit)

RYO (roll-your-own) synchronization:

- atomic_ref / atomic_def

Example: Monte Carlo Pi Calculation

- Estimate π using the “dartboard method” ($r=1$)
 - Area of square = $(2r)^2 = 4$
 - Area of circle = $\pi r^2 = \pi$
- Randomly throw darts at (x, y) positions inside the square
- If $x^2 + y^2 < 1$, then dart is inside the unit circle
- Calculate percentage that fall in the unit circle:
 - ratio = # darts inside circle / # darts inside square
- Compute π :
 - $\pi = 4 \times \text{ratio}$



$$A_S = (2r)^2 = 4r^2$$

$$A_C = \pi r^2$$

$$\pi = 4 \times \frac{A_C}{A_S}$$

```

program caf_pi
  implicit none

  integer      :: j
  integer      :: seed(2)
  integer*8    :: N_steps, i_step, hits
  double precision :: x, y
  double precision :: pi_sum, pi
  double precision :: pi_global[*]

  seed(1) = 17*this_image()
  call random_seed(put=seed)
  hits = 0_8
  N_steps = 10000000_8
  do i_step=1_8, N_steps
    call random_number(x)
    call random_number(y)
    if ( (x*x + y*y) <= 1.d0) then
      hits = hits + 1_8
    endif
  enddo

```

```

  pi_global = &
    4.d0*db1e(hits)/db1e(N_steps)

  SYNC ALL

  if (this_image() == 1) then

    pi_sum = 0.d0
    do j=1,num_images()
      pi_sum = pi_sum + pi_global[j]
    enddo

    pi = pi_sum / num_images()
    print *, 'pi = ', pi

  endif

end program caf_pi

```

Unified Parallel C (UPC)

- <https://upc-lang.org/>
- **UPC** (Unified Parallel C) is an extension to C, with the following constructs:
 - An explicitly parallel execution model (SPMD)
 - A shared address space
 - Synchronization primitives and a memory consistency model
 - Explicit communication primitives, e.g. `upc_memput`
 - Memory management primitives
- Multiple implementations:
 - Cray UPC
 - gcc version of UPC: <http://www.gccupc.org/>
 - Berkeley UPC: <http://upc.lbl.gov/>
- Most widely used on irregular / graph problems today

“Hello, world!” in UPC

```
/* needed for UPC extensions */
#include <upc.h>
#include <stdio.h>

int main() {
    printf("Hello from UPC thread %d out of %d:\n",
          MYTHREAD, THREADS);
    return 0;
}
```

- Any legal C program is also a legal UPC program
- If you compile and run it as UPC with P threads, it will run P copies of the program.

Note: some of the materials are borrowed from Kathy Yelick's presentation and Tarek El-Ghazawi's

- https://people.eecs.berkeley.edu/~demmel/cs267_Spr14/Lectures/lecture08-PGAS-yelick14_4pp.pdf
- http://upc.gwu.edu/downloads/upc_tut04.pdf

Compared to MPI “Hello, world!” in C

```
#include "mpi.h"
#include <stdio.h>

int main(int argc, char* argv[])
{
    int rank, size;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("Hello, world! I am process %d of %d\n", rank, size);
    MPI_Finalize();

    return 0;
}
```

Compared to OpenMP “Hello, world!” in C

```
#include <stdio.h>
#include <omp.h>
int main()
{
    int nthreads, tid;
    #pragma omp parallel private(tid)
    {
        tid = omp_get_thread_num();
        printf("Hello, world! I am thread %d\n", tid);
        #pragma omp barrier
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
    }
    return 0;
}
```


UPC on Cori

UPC is supported on Cori through 2 different implementations:

Berkeley UPC and Cray UPC

<http://www.nersc.gov/users/computational-systems/cori/programming/additional-programming-models/>

- **Berkeley UPC**

Berkeley UPC (BUPC) provides a portable UPC programming environment consisting of a **source translation front-end** (which in turn relies on a user-supplied C compiler underneath) and a runtime library based on **GASNet**. The latter is able to take advantage of advanced communications functionality of the Cray Aries interconnect on Cori, such as remote direct memory access (RDMA).

- **Cray UPC**

UPC is directly supported under Cray's compiler environment through their PGAS runtime library.

Man page: [intro_pgas\(7\)](#)

Berkeley UPC on Cori

- BUPC is available via the **bupc** module on Cori, which provides
 - The **upcc** compiler wrapper (all 3 programming environments, Intel, GNU & Cray, are supported by BUPC for use as the underlying C compiler)
 - The **upcrun** launcher wrapper (which initializes the environment and calls **srun**)

- Compiling and running UPC application with BUPC on Cori:

```
cori08> module load bupc
```

```
cori08> upcc upc_hello.c -o upc_hello.x
```

```
cori08> salloc -N 2 -t 10:00 -p debug -C haswell
```

```
nid00461> upcrun -n 64 ./upc_hello.x
```

Cray UPC on Cori

- To use Cray UPC on Cori, simply switch to the Cray compiler environment and supply the '**-h upc**' option when calling cc.
- Compiling and running UPC application with Cray UPC on Cori:

```
cori08> module swap PrgEnv-intel PrgEnv-cray
```

```
cori08> cc -h upc upc_hello.c -o upc_hello.x
```

```
cori08> salloc -N 2 -t 10:00 -p debug -C haswell
```

```
nid00461> ulimit -v unlimited # may not be necessary
```


```
nid00461> srun -n 64 ./upc_hello.x
```

UPC Execution Model

- A number of **threads** working independently in a SPMD fashion
 - UPC threads usually implemented as OS **processes**!
 - Number of threads specified at compile-time or run-time; available as program variable **THREADS**
 - **MYTHREAD** specifies thread index (0 .. **THREADS**-1)
- Synchronization when needed
 - Barriers
 - Locks
 - Memory consistency control

Private vs. Shared Variables in UPC

- Normal C variables and objects are allocated in the private memory space for each thread.
- Shared variables are allocated only once, with thread 0
`shared int ours; // use sparingly: performance`
`int mine;`
- Shared variables may not have dynamic lifetime: may not occur in a function definition, except as static.



Pi in UPC: Shared Memory Style

```
#include <stdio.h>
#include <math.h>
#include <upc.h>

int hit(){
    int const rand_max = 0xFFFFFFFF;
    double x = ((double) rand()) /
               RAND_MAX;
    double y = ((double) rand()) /
               RAND_MAX;
    if ((x*x + y*y) <= 1.0) {
        return(1);
    } else {
        return(0);
    }
}
```

```
shared int hits;

int main(int argc, char **argv) {
    int i, my_trials = 0;
    int trials = atoi(argv[1]);
    my_trials = (trials + THREADS - 1) /
               THREADS;
    srand(MYTHREAD*17);
    for (i=0; i < my_trials; i++)
        hits += hit();
    upc_barrier;
    if (MYTHREAD == 0)
        printf("PI estimated to %f.",
              4.0*hits/trials);
}
```

Race Condition!

Shared Arrays are Cyclic by default

- Shared scalars always live in thread 0
- Shared arrays are spread over the threads
- Shared array elements are spread across the threads

```
shared int x[THREADS]      /* 1 element per thread */  
shared int y[3][THREADS]   /* 3 elements per thread */  
shared int z[3][3]         /* 2 or 3 elements per thread */
```

- In the pictures below, assume THREADS = 4
 - Blue elements have affinity to thread 0

x 

y 

z 

z is not logically
blocked by columns!

UPC Shared Array Example

`shared double a[3][THREADS]`

<code>a[0][0]</code>
<code>a[1][0]</code>
<code>a[2][0]</code>

Thread 0

<code>a[0][1]</code>
<code>a[1][1]</code>
<code>a[2][1]</code>

Thread 1

<code>a[0][2]</code>
<code>a[1][2]</code>
<code>a[2][2]</code>

Thread 2

<code>a[0][3]</code>
<code>a[1][3]</code>
<code>a[2][3]</code>

Thread 3

Pi in UPC: Shared Array Version

```
shared int all_hits[THREADS];

int main(int argc, char **argv) {
    int i, hits, my_trials = 0;
    int trials = atoi(argv[1]);
    my_trials = (trials + THREADS - 1) / THREADS;
    srand(MYTHREAD*17);
    for (i=0; i < my_trials; i++)
        all_hits[MYTHREAD] += hit();

    upc_barrier;
    if (MYTHREAD == 0) {
        for (i=0; i < THREADS; i++)
            hits += all_hits[i];
        printf("PI estimated to %f.", 4.0*hits/trials);
    }
}
```

UPC Global Synchronization

- UPC has two basic forms of barriers:
 - Barrier: block until all other threads arrive
 - Split-phase barriers
- Optional labels allow for debugging

```
#define MERGE_BARRIER 12
if (MYTHREAD%2 == 0) {
    ...
    upc_barrier MERGE_BARRIER;
} else {
    ...
    upc_barrier MERGE_BARRIER;
}
```

Synchronization - Locks

- Locks in UPC are represented by an opaque type: **upc_lock_t**
- Locks must be allocated before use
 - Allocates 1 lock, pointer to all threads:
`upc_lock_t *upc_all_lock_alloc(void);`
 - Allocates 1 lock, pointer to one thread:
`upc_lock_t *upc_global_lock_alloc(void);`
- To use a lock (at start and end of critical region)
`void upc_lock(upc_lock_t *l);`
`void upc_unlock(upc_lock_t *l);`
- Locks can be freed when not in use
`void upc_lock_free(upc_lock_t *ptr);`

Pi in UPC: Shared Memory Style

```
shared int hits;
```

```
int main(int argc, char **argv) {  
    int i, my_hits, my_trials = 0;  
    upc_lock_t *hit_lock = upc_all_lock_alloc();  
    int trials = atoi(argv[1]);  
    my_trials = (trials + THREADS - 1) / THREADS;  
    srand(MYTHREAD*17);  
    for (i=0; i < my_trials; i++) my_hits += hit();
```

```
    upc_lock(hit_lock);  
    hits += my_hits;  
    upc_unlock(hit_lock);
```

accumulate across threads

```
    upc_barrier;  
    if (MYTHREAD == 0)  
        printf("PI estimated to %f.", 4.0*hits/trials);  
}
```

UPC (Value-Based) Collectives

- A portable library of collectives on scalar values (not arrays)

Example: `x = bupc_allv_reduce(double, x, 0, UPC_ADD);`

`TYPE bupc_allv_reduce(TYPE, TYPE value, int root, upc_op_t op);`

- 'TYPE' is the type of value being collected
- root is the thread ID for the root (e.g., the source of a broadcast)
- 'value' is both the input and output (must be a “variable” or l-value)
- op is the operation: UPC_ADD, UPC_MULT, UPC_MIN, ...
- Computational Collectives: reductions and scan
- Data Movement Collectives: broadcast, scatter, gather
- Portable implementation available from:
 - http://upc.lbl.gov/download/dist/upcr_preinclude/bupc_collectivev.h
- UPC also has more general collectives over arrays
 - <https://upc-lang.org/assets/Uploads/spec/upc-lang-spec-1.3.pdf>

Pi in UPC: Data Parallel Style

```
#include <bupc_collectivev.h>
shared int hits;

int main(int argc, char **argv) {
    int i, my_hits, my_trials = 0;
    int trials = atoi(argv[1]);
    my_trials = (trials + THREADS - 1) / THREADS;
    srand(MYTHREAD*17);
    for (i=0; i < my_trials; i++) my_hits += hit();

    //                                type, input, thread, op
    bupc_allv_reduce(int, my_hits, 0, UPC_ADD);
    // upc_barrier;

    if (MYTHREAD == 0)
        printf("PI estimated to %f.", 4.0*my_hits/trials);
}
```

UPC++



- <https://bitbucket.org/upcxx/upcxx/wiki/Home>
- UPC++ paper:
<https://web.eecs.umich.edu/~akamil/papers/ipdps14.pdf>
- UPC++ is a PGAS extension for C++
 - “Compiler-free” approach using C++ templates and runtime libraries
- UPC++ provides significantly more expressiveness than UPC:
 - an object-oriented PGAS programming model in the context of C++
 - useful parallel programming idioms unavailable in UPC, such as **asynchronous remote function invocation** and **multidimensional arrays**
 - an easy on-ramp to PGAS programming through interoperability with other existing parallel programming systems (e.g., MPI, OpenMP, CUDA)
- UPC++ performance is close to UPC!

“Hello, world!” in UPC++

```
// upcxx_hello.cpp - Hello world in UPC++
#include <upcxx.h>
#include <iostream>

using namespace upcxx;

int main (int argc, char **argv)
{
    init(&argc, &argv);
    std::cout << "I'm rank " << myrank() << " of " << ranks() << "\n";
    finalize();
    return 0;
}
```


UPC++ on Cori

- UPC++ is available via the **upcxx** module on Cori, which provides the **upc++** compiler wrapper

- Compiling and running UPC++ application on Cori:

```
cori08> module load upcxx
```

```
cori08> upc++ -std=c++11 upcxx_hello.cpp -o upcxx_hello.x
```

```
cori08> salloc -N 2 -t 10:00 -p debug -C haswell
```

```
nid00461> srun -n 64 ./upcxx_hello.x
```

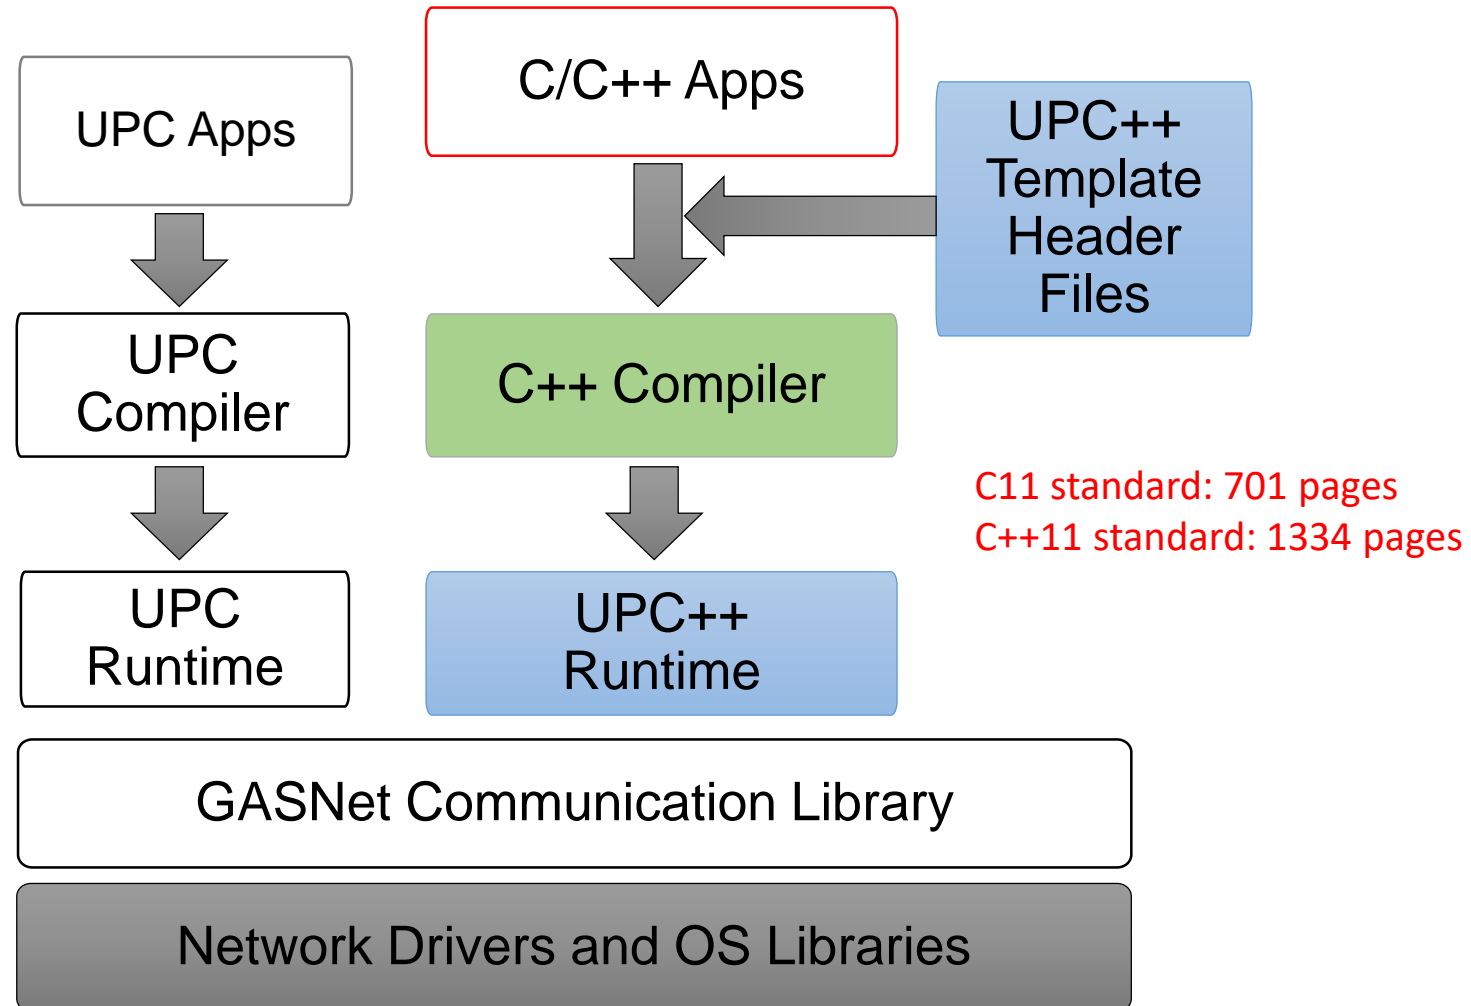
1. Compiling UPC++ Applications:

<https://bitbucket.org/upcxx/upcxx/wiki/Compiling%20UPC++%20Applications>

2. Running UPC++ Applications:

<https://bitbucket.org/upcxx/upcxx/wiki/Running%20UPC++%20Applications>

UPC++ Software Stack



UPC++ Introduction

- Shared variable

```
shared_var<int> s;    // int in the shared space
```

- Global pointers (to remote data)

```
global_ptr<LLNode> g;    // pointer to shared space
```

- Shared arrays

```
shared_array<int> sa(8);    // array in shared space
```

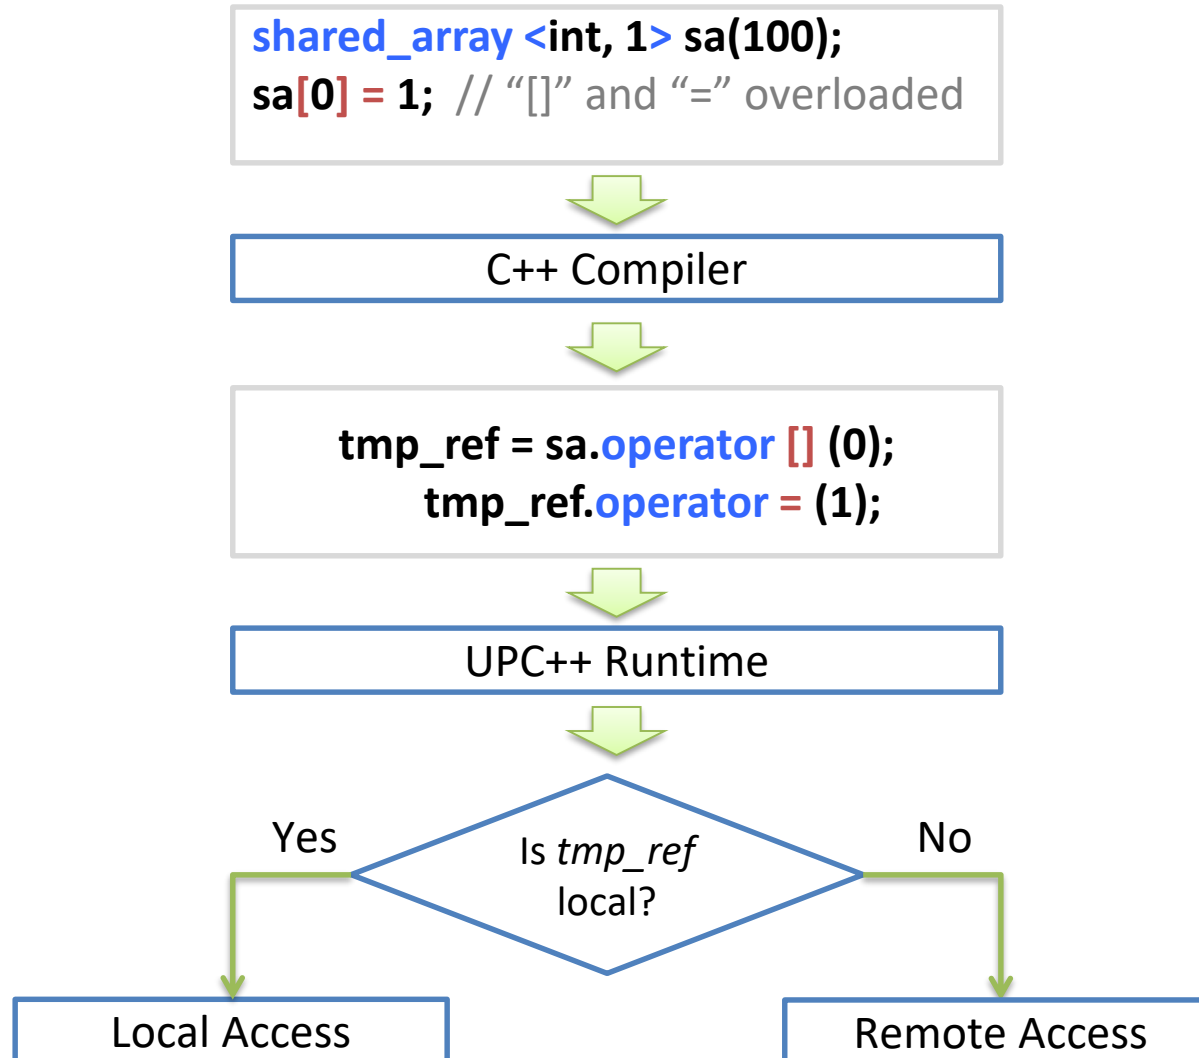
- Locks

```
shared_lock l;    // lock in shared space
```

- Default execution model is SPMD, but with optional async

```
async(place)(Function f, T1 arg1,...);  
wait();    // other side does poll()
```

UPC++ Translation Example



Dynamic Global Memory Management

- Global address space pointers (pointer-to-shared)

```
global_ptr<data_type> ptr;
```

- Dynamic shared memory allocation

```
global_ptr<T> allocate<T>(uint32_t where, size_t count);  
void deallocate(global_ptr<T> ptr);
```

- Example: allocate space for 512 integers on rank 2

```
global_ptr<int> p = allocate<int>(2, 512);
```

One-Sided Data Transfer Functions

```
// Copy count elements of T from src to dst  
upcxx::copy<T>(global_ptr<T> src, global_ptr<T> dst, size_t count);
```

```
// Non-blocking version of copy  
upcxx::async_copy<T>(global_ptr<T> src, global_ptr<T> dst,  
                      size_t count);
```

```
// Synchronize all previous asyncs  
upcxx::async_wait();
```

Asynchronous Task Execution

- C++ 11 **async** function

```
std::future<T> handle = std::async(Function&& f, Args&&... args);  
handle.wait();
```

- UPC++ **async** function

```
// Remote Procedure Call  
upcxx::async(place)(Function f, T1 arg1, T2 arg2,...);  
upcxx::wait();
```

```
// Explicit task synchronization  
upcxx::event e;  
upcxx::async(place, &e)(Function f, T1 arg1, ...);  
e.wait();
```

Further Readings

- Coarrays in the next Fortran Standard:
<ftp://ftp.nag.co.uk/sc22wg5/N1801-N1850/N1824.pdf>
- OpenSHMEM Specification 1.3:
http://www.openshmem.org/site/sites/default/site_files/OpenSHMEM-1.3.pdf
- OpenSHMEM Tutorials: <http://openshmem.org/site/Documentation/Tutorials>
- UPC Language Specifications, v1.3:
<https://upc-lang.org/assets/Uploads/spec/upc-lang-spec-1.3.pdf>
- Berkeley UPC User's Guide: <http://upc.lbl.gov/docs/user/index.shtml>
- PGAS with HPC, by Kathy Yelick
- Programming in UPC, by Tarek El-Ghazawi
- GASNet: <http://gasnet.lbl.gov/>
- A tutorial of UPC++: <https://bitbucket.org/upcxx/upcxx/wiki/Tutorial>
- UPC++, by Yili Zheng

Thank you!